

Recentered influence functions (RIFs) in Stata: RIF regression and RIF decomposition

Fernando Rios-Avila
Levy Economics Institute of Bard College
Annandale-on-Hudson, NY
friosavi@levy.org

Abstract. Recentered influence functions (RIFs) are statistical tools popularized by Firpo, Fortin, and Lemieux (2009, *Econometrica* 77: 953–973) for analyzing unconditional partial effects on quantiles in a regression analysis framework (unconditional quantile regressions). The flexibility and simplicity of these tools have opened the possibility to extend the analysis to other distributional statistics using linear regressions or decomposition approaches. In this article, I introduce one function and two commands to facilitate the use of RIFs in the analysis of outcome distributions: `rifvar()` is an `egen` extension used to create RIFs for a large set of distributional statistics, `rifhdreg` facilitates the estimation of RIF regressions enabling the use of high-dimensional fixed effects, and `oaxaca_rif` implements Oaxaca–Blinder decomposition analysis (RIF decompositions).

Keywords: `st0588`, `rifvar()`, `rifhdreg`, `rifsureg2`, `oaxaca_rif`, `uqreg`, recentered influence functions, unconditional partial effects, unconditional quantile regression, RIF regressions, distributional statistics, Oaxaca–Blinder, RIF decomposition

1 Introduction

Influence functions (IFs) are statistical tools that have been used for analyzing the robustness of distributional statistics, or functionals, to small disturbances in data (Cowell and Flachaire 2007) or as a simplified strategy to estimate asymptotic variances of complex statistics (Cowell and Flachaire 2015; Deville 1999). More recently, Firpo, Fortin, and Lemieux (2009) suggested the use of IFs—specifically, recentered influence functions (RIFs)—as tools to analyze the impact that changes in the distribution of explanatory variables, X , has on the unconditional distribution of Y .

The method introduced by Firpo, Fortin, and Lemieux (2009) focused on the estimation of unconditional quantile regression (UQR), which allows the researcher to obtain partial effects of explanatory variables on any unconditional quantile of the dependent variable. The simplest version of this methodology, referred to as RIF–OLS (ordinary least squares), is easily implemented using the community-contributed command `rifreg` by the same authors. As part of their conclusions, the authors highlight the potential extensions of this strategy for analyzing other distributional statistics, as well as the

potential usefulness for generalizing the traditional Oaxaca–Blinder (OB) decomposition for analyzing differences of outcome distributions across groups.¹

After its introduction, UQR became a popular method for analyzing and identifying the distributional effects on outcomes in terms of changes in observed characteristics in areas such as labor economics, income and inequality, health economics, and public policy. The potential simplicity and flexibility the methodology offers for the analysis of any distributional statistics also motivated subsequent research to expand the use of RIFs in the framework of regression analysis.

In a recently published article, Firpo, Fortin, and Lemieux (2018) discuss the application of RIF regressions for the variance and Gini coefficients, with emphasis on the generalization of the OB decomposition using RIFs.² Borgen (2016), building on the work of Firpo, Fortin, and Lemieux (2009), provides the command `xtrifreg` for the efficient estimation of UQR in the presence of a single high-dimensional fixed effect, but `xtrifreg` cannot be used for other distributional statistics.

Cowell and Flachaire (2007), for their analysis on the sensitivity of inequality measures to the presence of extreme values, provide IFs for the most commonly used inequality indices, including the Atkinson index, generalized entropy index, and logarithmic variance index. Essama-Nssah and Lambert (2012), who discuss the use of RIF regressions and OB decompositions for the analysis of distributional changes, provide a large set of IFs and RIFs for distributional statistics relevant for policy analysis, including Lorenz and generalized Lorenz ordinates; Foster–Greer–Thorbecke (FGT) poverty indices; and Watts and Sen poverty indices. Most recently, Heckley, Gerdtham, and Kjellsson (2016) examine the use of RIFs for measures of health inequality, with emphases on bivariate rank-dependent concentration indices.

While RIF regressions and RIF decompositions have become important analytical tools in the empirical literature, there are only limited attempts to provide a simplified framework to allow the use of RIFs as a standard analytical tool. Within Stata, only the community-contributed command `xtrifreg` is readily available from the *Stata Journal* archives. The command `rifireg` is also available from the author’s website, but it is limited to the analysis of rank-dependent indices.³ The command that started it all, `rifreg`, is limited in the estimation of RIF statistics and is not available in the Statistical Software Components archives, although it can be accessed manually from the author’s website. Furthermore, while RIFs and reweighted regressions are broadly

-
1. The use of RIF regressions within the OB decomposition approach has been discussed in the review of decomposition methods by Fortin, Lemieux, and Firpo (2011) and more recently in Firpo, Fortin, and Lemieux (2018).
 2. While recently published, the working version of the article dates from 2007 and was cited in their 2009 article as part of the extensions on the use of RIFs in regression analysis. The command `rifreg` estimates all the RIFs proposed in the article.
 3. The command `rifireg` can be obtained from <https://sites.google.com/site/gawainheckley/home/stata-code>.

used for the generalization of the OB decomposition for statistics beyond the mean, existing commands can be used only to obtain decompositions based on a multistep process.⁴

In this article, I introduce one function and two commands that aim to facilitate the use of RIFs for regressions and decomposition analysis. The first function, `rifvar()`, is a byable plugin extension that works with `egen` and can be used to estimate the RIF for a large set of distributional statistics, such as those described in Firpo, Fortin, and Lemieux (2018); Essama-Nssah and Lambert (2012); Cowell and Flachaire (2007); and Heckley, Gerdtham, and Kjellsson (2016). The first command, `rifhdreg`, is a wrapper command for `regress` and `reghdfe` (Correia 2016) that in combination with `rifvar()` is used to estimate RIF regressions in the presence of high-dimensional fixed effects as well as inequality treatment effects. The second command, `oaxaca_rif`, is a wrapper around `oaxaca` (Jann 2008) that can be used to implement standard and reweighted OB decompositions (see Fortin, Lemieux, and Firpo [2011] and Firpo, Fortin, and Lemieux [2018]), implementing the multistep procedure described in Firpo, Fortin, and Lemieux (2018).

The rest of this article is structured as follows. Section 2 overviews what IFs are and how they are estimated. Section 3 introduces and explains the use of `rifvar()` to estimate RIF variables. Section 4 describes the use of `rifhdreg` for the estimation of standard RIF regressions, as well as for the estimation of inequality treatment effects. Section 5 describes the use of `oaxaca_rif` for the estimation of standard and reweighted decompositions using RIF decomposition. Section 6 concludes.

2 RIF and distributional statistics

2.1 Distributional statistics: Basics

When analyzing social welfare, inequality, poverty, or any other measure that describes the distributional characteristics of an outcome of interest, it is necessary to have access to one of the following pieces of information. The most common scenario is to have access to the full set of values corresponding to each observation in the population or sample. If the population or sample is of finite size n , it can be referred to as $Y = [y_1, y_2, \dots, y_n]$, where each y_i is the outcome of interest (for example, income) of the i th person.

The second scenario is where we know the distribution of the outcome, based on either the cumulative distribution function (c.d.f.) or the probability distribution function (p.d.f.). Using the function $F_Y(\cdot)$ to refer to the c.d.f. and $f_Y(\cdot)$ to the p.d.f., the vector of information required for analyzing distributions can be more briefly written as a set of ordered pairs, $F_Y = [\{y, F_Y(y)\} | y \in \mathbb{R}]$ or $f_Y = [\{y, f_Y(y)\} | y \in \mathbb{R}]$, where y

4. For example, Firpo, Fortin, and Lemieux (2018) describe a three-step process for the implementation of the RIF-reweighted decomposition: 1) estimation of reweighting factors using probit or logit models, 2) estimation of the corresponding RIFs using `rifreg`, and 3) application of the OB decomposition using the `oaxaca` command.

represents any real number (generally positive when referring to income).⁵ This simply means that if one has access to any of these vectors of information (Y, F_Y, f_y) , any distributional statistic can be derived.

Let us call $v(\cdot)$ a functional that uses all the information contained in Y , F_Y , or f_Y to estimate a distributional statistic of Y . This functional can be used to estimate statistics relevant to policy analysis, such as the mean, q th quantile, poverty indices, or inequality indices. To measure the impact a change in the distribution of income will have on the distributional statistic, one can simply compare the indices by swapping the c.d.f. from the observed distribution F_Y to the ex-post distribution G_Y .⁶ This can be written as follows:

$$\Delta v = v(G_Y) - v(F_Y) \quad (1)$$

Thus, Δv is the change in the distributional statistic generated by the change in distribution from $F_y \rightarrow G_y$. This change can be as large as implying that everyone in the population receives a fixed transfer (shifting the function $F_y(\cdot)$ to the right)⁷ or as simple as having a new person (with random income) added to the sample, changing the rankings of everyone in the sample. The first scenario is a simplified example of what DiNardo, Fortin, and Lemieux (1996) used for analyzing changes in the distribution of wages. The second scenario is an exercise that can be used for understanding the definition of IFs and RIFs.

2.2 Estimations of IFs and RIFs: Gâteaux derivative

The thought experiment of adding a new person to a sample can also be considered as a case of data contamination in the original distribution, and (1) can be used to estimate the influence of this thought experiment on the statistic v . Standardizing the change in the statistic Δv with respect to some measure that quantifies the change of the distribution $[\Delta(G_Y - F_Y)]$ provides a quantification of the rate of change of the statistic v associated with the change of the distribution of Y from $F_Y \rightarrow G_Y$.

$$\Delta^s v = \frac{\Delta v}{\Delta(G_Y - F_Y)} = \frac{v(G_Y) - v(F_Y)}{\Delta(G_Y - F_Y)}$$

5. Note that the functions $F_Y(\cdot)$ and $f_Y(\cdot)$ are not arbitrary and obey a strict set of properties and relations among them to represent well-defined distribution functions: $f_Y(y) \geq 0 \forall h \in \mathbb{R}$, $\int_{-\infty}^y f_Y(x)dx = F_Y(y)$, $dF_Y(y) = f_Y(y)$, $F_Y(-\infty) = 0$, $F_Y(\infty) = 1$, and $F_Y(y_1) \leq F_Y(y_2) \Leftrightarrow y_1 \leq y_2$.

6. $G_Y(\cdot)$ and its counterpart $g_Y(\cdot)$ have the same properties as $F_Y(\cdot)$ and $f_Y(\cdot)$, respectively.

7. This thought experiment can also be as if people with low income suddenly disappear whereas people with higher income become more numerous.

This process can be extended to measure $\Delta^s v$ for an infinitesimally small change in the distribution function from $F_Y \rightarrow G_Y$. This idea is what lies behind the Gâteaux derivative, which is a generalization of the directional derivative of a functional. The derivative is used to construct IFs, which can be used as measures of robustness of functionals to data outliers (Hampel 1974) and which facilitate the visualization of the structure of the distributional statistic as a function of the available data. Before we proceed to the formal definition of the IF, it would be useful to formalize the thought experiment just described.

Assume that the observation to be introduced in the sample has an income equal to y_i . Because this is the only element of that distribution, its c.d.f., $H_{y_i}(y)$, can be characterized as follows:

$$\begin{aligned} H_{y_i}(y) &= 0 \quad \forall y < y_i \\ H_{y_i}(y) &= 1 \quad \forall y \geq y_i \end{aligned}$$

This indicates that the distribution H_{y_i} puts mass only at the value y_i .⁸ With this definition, the distribution G_Y can be redefined as a combination of the distributions F_Y and H_{y_i} :

$$G_Y = (1 - \varepsilon)F_Y + \varepsilon H_{y_i} \quad (2)$$

ε is strictly smaller than 1 but larger than 0. In other words, G_Y is the resulting distribution when the original distribution F_Y changes in the direction of H_{y_i} . This transformation can also be thought of as a contamination or perturbation of the distribution F_Y in the direction of H_{y_i} . Equation (2) also helps quantify the magnitude of the change in the distribution when moving from $F_Y \rightarrow G_Y$ (simply, ε). Figure 1 provides a graphical example of the changes we would observe as a result of this contamination of the distribution function F_Y .

8. Conversely, this means that $dH_{y_i}(y) = 0 \quad \forall y \neq y_i$ and $dH_{y_i}(h) = \infty$ if $y = y_i$.




Figure 1. Comparison between original and contaminated distributions. NOTE: The figure represents the contamination of the original distribution using $\varepsilon = 0.2$.

With this last concept in place, we can finally provide the formal definition of the IF:

$$\text{IF} \{y_i, v(F_Y)\} = \lim_{\varepsilon \rightarrow 0} \frac{v \{(1 - \varepsilon)F_Y + \varepsilon H_{y_i}\} - v(F_Y)}{\varepsilon} = \frac{\partial v(F_Y \rightarrow H_{y_i})}{\partial \varepsilon}$$

The IF is a directional derivative that shows the rate of change of the distributional statistic v caused by an infinitesimally small change in the distribution F_y in the direction of H_{y_i} . It can also be interpreted as the influence that observation y_i has on the estimation of the distributional statistic v . Take the distributional statistic mean as an example. The statistic and its IF are defined as follows:

$$\mu_{F_y} = \int y dF_y$$

$$\text{IF} \{y_i, v(F_y)\} = y_i - \mu_y$$

The IF indicates that if the distribution F_y were to be contaminated, adding more weight to observations with income y_i , the mean would change at a rate of $y_i - \mu_y$. If the change in the distribution is ε , the change of the mean would be $\Delta \mu_y = \varepsilon \times \text{IF} \{y_i, v(F_Y)\}$. Note that the IF will be different for each point of reference y_i (contamination point) used for its estimation.

Instead of using the IF directly, Firpo, Fortin, and Lemieux (2009) propose using the recentered version of the statistics, the RIF, which is equivalent to the first two terms of the von Mises (1947) linear approximation of the corresponding distributional statistic v :

$$\text{RIF} \{y_i, v(F_Y)\} = v(F_Y) + \text{IF} \{y_i, v(F_Y)\} \quad (3)$$

This expression can be loosely interpreted as the relative contribution that observation y_i has on the construction of the statistic v . It can also be interpreted as an

approximation of the statistic v after considering the influence of observation y_i . Returning to the example of the mean, the $\text{RIF}(y_i, \mu_{F_y})$ is defined as the observation itself, y_i , and it is easy to see that y_i is the relative contribution that observation i has on the construction of the mean.

As discussed in Firpo, Fortin, and Lemieux (2009); Cowell and Flachaire (2015); and Essama-Nssah and Lambert (2012), the IF and the RIF have the following properties:

$$\int \text{IF} \{y, v(F_Y)\} dF_Y = 0 \quad (4)$$

$$\int \text{RIF} \{y, v(F_Y)\} dF_Y = v(F_Y) \quad (5)$$

$$v(F_y) \sim N \left\{ v(F_Y), \frac{\sigma_{\text{IF}}^2}{N} \right\} \quad (6)$$

$$\sigma_{\text{IF}}^2 = \int \text{IF} \{y, v(F_Y)\}^2 dF_Y \quad (7)$$

Equation (4) implies that the expected value of IFs constructed using all values of y_i is equal to 0. Because of this property, (5) implies that the expected value of RIFs is equal to the distributional statistic itself. Equations (6) and (7) indicate that the asymptotic variance of any statistic can be obtained by estimating the variance of the IF or the RIF based on sample data.

As shown in (3), choosing between IFs or RIFs as dependent variables has no impact in terms of regression analysis other than changes in the intercept in the case of RIF-OLS. This happens because IFs and RIFs are equivalent to each other up to a constant. The advantage of using RIFs, as shown in (5), is that one can use them to recover the underlying distributional statistics using simple averages. This property facilitates the interpretation of RIF regressions and enables the implementation of decomposition analysis.

3 Estimating RIFs: `rifvar()` egen function

The estimation of RIFs is a task of variable complexity. Some statistics have simple mathematical expressions that require few lines of code to define the corresponding RIF. The easiest example is the RIF for the mean because the RIF mean for any value y_i is simply itself. However, other statistics may require many intermediate steps to correctly define their corresponding RIFs.

The community-contributed command `rifreg` can be used to estimate RIF regressions and create their corresponding RIFs, but it is limited to the analyses of the variance, quantile, and Gini coefficient. If you are interested in analyzing other distributional

statistics, I suggest the use of a new function called `rifvar()`.⁹ This is a plugin command that adds new functions to the command `egen`, facilitating the estimation of RIFs for a large set of distributional statistics.¹⁰

The syntax of the command is as follows:

```
egen [ type ] newvar = rifvar(varname) [ if ] [ in ] [ , by(varname)  
    weight(varname) seed(str) RIF_options ]
```

where *varname* is the variable being analyzed and *newvar* is the new variable name where the RIF will be stored, given the restrictions set by *if*, *in*, or both. All statistics allow for use of the option *by*(*varname*), used to indicate the variables over which the RIF will be estimated (that is, sex or race group), and the option **weight**(*varname*), used to indicate weights for the estimation of the RIFs. **seed**(*str*) indicates a particular seed for replication of rank-dependent indices.

RIF_options allow the user to specify which distributional statistics are used to obtain the RIF statistic. Table 1 provides a detailed list of the statistics that are currently available for estimation. In appendix B, I provide a summary of performance simulation to show how well the RIF standard errors approximate the simulated standard errors for all of these statistics.

Table 1. `rifvar()` *RIF_options*

Option	Description
mean	sample mean
var	variance
q (<i>#p</i>) [kernel (<i>kernel</i>) bw (<i>#</i>)]	<i>p</i> th quantile, where $0 < \#p < 100$; options kernel () and bw () are optional; default is kernel (gaussian) kernel; all kernel functions available for command kdensity are also allowed; † default is Silverman's plug-in optimal bandwidth*
iqr (<i>#p1</i> <i>#p2</i>) [kernel (<i>kernel</i>) bw (<i>#</i>)]	interquantile range: $q(\#p2) - q(\#p1)$, where $0 < p_1 \leq p_2 < 100$; bw () and kernel () options are the same as for the quantile case
gini	Gini inequality index
cvar	coefficient of variation
std	standard deviation
iqratio (<i>#p1</i> <i>#p2</i>) [kernel (<i>kernel</i>) bw (<i>#</i>)]	interquantile ratio: $q(\#p2)/q(\#p1)$, where $0 < p_1 \leq p_2 < 100$; bw () and kernel () options are the same as for the quantile case

Continued on next page

9. Internally, the function is stored in a file named `_grifvar.ado`. `rifvar()` builds on the `rifreg` command and the do-file provided in the appendix in Heckley, Gerdtham, and Kjellsson (2016). All codes were adapted to allow for estimations by groups.

10. See appendix A for the full set of distributional statistics, formulas, and sources.

Option	Description
<code>entropy(#a)</code>	generalized entropy index with sensitivity parameter <code>#a</code>
<code>atkin(#e)</code>	Atkinson inequality index with inequality aversion <code>#e > 0</code>
<code>logvar</code>	logarithmic variance (different from variance of logarithms)
<code>glor(#p)</code>	generalized Lorenz ordinate at <code>#p</code> , where $0 < #p < 100$
<code>lor(#p)</code>	Lorenz ordinate at <code>#p</code> , where $0 < #p < 100$
<code>ucs(#p)</code>	share of income held by richest $1 - p\%$; $1 - \text{lor}(\#p)$
<code>iqsr(#p1 #p2)</code>	interquantile share ratio: $(1 - \text{lor}(\#p2)) / \text{lor}(\#p1)$, where $0 < \#p1 \leq \#p2 < 100$
<code>mcs(#p1 #p2)</code>	share of income held by people between <code>#p1</code> and <code>#p2</code> : $\text{lor}(\#p2) - \text{lor}(\#p1)$, where $0 < \#p1 < \#p2 < 100$
<code>pov(#a) pline(# varname)</code>	FGT poverty index given sensitivity parameter <code>#a</code> ; <code>pov(0)</code> returns the poverty head count, <code>pov(1)</code> returns the poverty gap, and <code>pov(2)</code> returns the poverty severity; FGT are defined based on the poverty line <code>pline()</code> , which can be a scalar (fixed poverty line) or a variable (variable poverty line)
<code>watts(#povline)</code>	Watts poverty index; requires a number or variable to define the poverty line
<code>sen(#povline)</code>	Sen poverty index; requires a number to define the poverty line
<code>tip(#p) pline(#)</code>	three I's of poverty (TIP) curve ordinate at <code>#p</code> for poverty line defined by <code>pline(#)</code> , where $0 < \#p < 100$
<code>agini</code>	absolute Gini
<code>acindex(varname)</code>	absolute concentration index using <code>varname</code> as the rank variable
<code>cindex(varname)</code>	concentration index using <code>varname</code> as the rank variable
<code>eindex(varname) lb(#) ub(#)</code>	Erreygers's index using <code>varname</code> as the rank variable, with lower bound <code>lb()</code> and upper bound <code>ub()</code> and where $\text{lb}() < \text{ub}()$
<code>arcindex(varname) lb(#)</code>	attainment relative concentration index using <code>varname</code> as the rank variable, with lower bound <code>lb()</code>
<code>srindex(varname) ub(#)</code>	shortfall relative concentration index using <code>varname</code> as the rank variable, with upper bound <code>ub()</code>
<code>windex(varname) lb(#) ub(#)</code>	Wagstaff concentration index using <code>varname</code> as the rank variable, with lower bound <code>lb()</code> and upper bound <code>ub()</code> and where $\text{lb}() < \text{ub}()$

Continued on next page

Option	Description
<code>rifown(str)</code> <code>rifopt(str)</code>	these options allow the use of <code>rifvar()</code> , <code>rifhdreg</code> , and <code>oaxaca_rif</code> with other community-contributed RIFs for statistics that are not available in the current list; this may include the estimation of RIFs using numerical methods such as jackknife methods (Efron 1982) or analytical methods not yet programmed; ¹¹ to use this, the new function should be implemented as an <code>egen</code> extension allowing for <code>weight()</code> and <code>by()</code> options, <code>rifown()</code> should contain the extension command name, and <code>rifopt()</code> should contain other options associated with the command; see file <code>_ghvarp().ado</code> as an example

NOTES: † To use other kernel functions, the full kernel name should be used and not its abbreviation (for example, `biweight` instead of `bi`), with one exception: to request using the Epanechnikov kernel, one should use `epan`.

* In contrast, `rifreg` uses the default bandwidth from `kdensity`. This accounts for small differences in the estimation of unconditional quantile RIFs compared with `rifreg`. See the `reghdfe` help file for an example replicating `rifreg` results.

4 RIF regression: `rifhdreg`

4.1 Standard RIF regression

As previously indicated, IFs and RIFs have been used in statistics as tools for analyzing the robustness of statistics to outliers and as methods to draw statistical inferences from complex statistics (Cowell and Flachaire 2015; Deville 1999; Efron 1982; Hampel 1974). A recently popularized use by Firpo, Fortin, and Lemieux (2009); Heckley, Gerdtham, and Kjellsson (2016); and Essama-Nssah and Lambert (2012) is the estimation of RIF regressions.

Firpo, Fortin, and Lemieux (2009) use this strategy to estimate unconditional partial effects (UPE) of small changes in the distribution of the independent variables \mathbf{X} on the distributional statistic v .¹² The authors use this strategy for the estimation of UQRs, using linear models as the easiest method for approximating these partial effects.¹³ Paraphrasing the original article, the estimation of RIF regressions can be described as follows.

11. I thank an anonymous referee for suggesting this addition to the command.

12. Firpo, Fortin, and Lemieux (2009, 958) call this an unconditional partial effect.

13. For the special case of quantiles, the RIF is defined as $q_Y(p) + [p - 1\{y \leq q_Y(p)\}]/[f\{q_Y(p)\}]$. Because the only element of this expression that varies across observations is $1\{y \leq q_Y(p)\}$, Firpo, Fortin, and Lemieux (2009) propose to model this using three methods: a linear probability model (RIF-OLS), a probit model (RIF-probit), or a nonparametric binomial model (RIF-NP). See appendix C for an additional command, `uqreg`, that allows the estimation of UQRs using binomial models such as probit or logit.

Assume there is a joint distribution function $dF_{Y,X}(y, \mathbf{x}) = dF_{Y|X}(y|\mathbf{X} = \mathbf{x})dF_X(\mathbf{x})$ that determines the linear and nonlinear relationships between the dependent variable Y and all independent or exogenous variables \mathbf{X} . Using this function, the p.d.f. and c.d.f. of the unconditional distribution of Y can be written as

$$dF_Y(y_i) = \int dF_{Y|X}(y_i|\mathbf{X} = \mathbf{x})dF_X(\mathbf{x}) \quad (8)$$

$$F_Y(y_i) = \int F_{Y|X}(y_i|\mathbf{X} = \mathbf{x})dF_X(\mathbf{x}) \quad (9)$$

These expressions simply state that the unconditional c.d.f. (and p.d.f.) of Y can be obtained by integrating the conditional distributions $F_{Y|X}(\cdot)$ (averaging) across all the realizations of \mathbf{X} . It also implies that if one assumes $F_{Y|X}(y|\mathbf{x})$ remains constant, then changes in the distribution of \mathbf{X} , say, from $F_X \rightarrow G_X$, will translate in changes of the unconditional distribution of Y .

$$G_Y = \int F_{Y|X}(y_i|\mathbf{X} = \mathbf{x})dG_X(\mathbf{x}) \quad (10)$$

Next assume that one is interested in analyzing the distributional statistic $v(F_Y)$. Based on (8) and (5), respectively, the statistic v can be rewritten as

$$v(F_Y) = \int \int \text{RIF}\{y, v(F_Y)\} dF_{y|X}(y|\mathbf{X} = \mathbf{x})dF_X(\mathbf{x}) \quad (11)$$

$$v(F_Y) = \int E[\text{RIF}\{y, v(F_Y)\}|\mathbf{X} = \mathbf{x}]dF_X(\mathbf{x}) \quad (12)$$

An intuitive interpretation of (11) is that any distributional statistics v can be written as the average of the RIFs weighted by the joint distribution $dF_{Y,X}(y, \mathbf{x}) = dF_{y|X}(y|\mathbf{X} = \mathbf{x})dF_X(\mathbf{x})$. Alternatively, (12) indicates that any distributional statistics can be obtained by averaging the conditional expectations of the RIFs with respect to \mathbf{X} across the distribution $dF_X(\mathbf{x})$. Equations (10) and (11) can also be used to state that if the distribution of the exogenous characteristics \mathbf{X} changes, say, from $F_X \rightarrow G_X$, assuming that the conditional distribution $F_{y|X}$ is constant, there will be a change in the unconditional distribution of Y from $F_Y \rightarrow G_Y$ that will translate into a change in the distributional statistic $v(F_Y) \rightarrow v(G_Y)$. Equation (12), which is used by Firpo, Fortin, and Lemieux (2009) to validate the use of RIFs in regression analysis, further suggests that the change in the distributional statistic caused by the change $F_X \rightarrow G_X$ can be approximated by

$$v(G_Y) - v(F_Y) = \int E[\text{RIF}\{y, v(F_Y)\}|\mathbf{X} = \mathbf{x}]d(G_X - F_X)(\mathbf{x}) \quad (13)$$

Equations (12) and (13) also imply that if $E[\text{RIF}\{y, v(F_Y)\}|\mathbf{X} = \mathbf{x}]$ can be modeled as a function of the distribution of \mathbf{X} , those results can be used to estimate how changes in the distribution of \mathbf{X} may affect the distributional statistic v .

Following Firpo, Fortin, and Lemieux (2009), the simplest approach to estimate RIF regressions is to assume a linear relationship between $\text{RIF}\{y, v(F_Y)\}$ and the explanatory variables \mathbf{X} . Under this assumption, OLS can be used to fit a linear model to capture how small changes in the distribution of the independent variables \mathbf{X} affect $v(F_Y)$. The difference with the standard OLS model is that RIF-OLS uses the $\text{RIF}\{y, v(F_Y)\}$ for each observation y_i in the data as the dependent variable and regresses it against all the variables of interest:¹⁴

$$\text{RIF}\{y, v(F_Y)\} = \mathbf{X}'\boldsymbol{\beta} + \varepsilon_i, \quad E(\varepsilon_i) = 0 \quad (14)$$

While the use of OLS directly relates the RIF regression to standard regression analysis, some differences in the interpretation exist. In the standard OLS, the typical interpretation of the coefficients is that a one-unit increase in \mathbf{X} will cause y to increase in $\boldsymbol{\beta}$ units (on average), everything else held constant. The interpretation from the RIF-OLS is slightly different. To obtain the UPE on the statistic v , one first needs to obtain unconditional expectations on both sides of (14):

$$v(F_Y) = E[\text{RIF}\{y, v(F_Y)\}] = E(\mathbf{X}'\boldsymbol{\beta}) + E(\varepsilon_i) = \bar{\mathbf{X}}'\boldsymbol{\beta} \quad (15)$$

$\bar{\mathbf{X}}$ is the unconditional mean of \mathbf{X} . From here, the UPE is given by

$$\frac{\partial v(F_Y)}{\partial \bar{X}_k} = \beta_k \quad (16)$$

Based on (16), the correct interpretation of the UPE is that if the distribution of x_k changes such that its unconditional average increases by one unit ($\Delta \bar{X}_k = 1$),¹⁵ the distributional statistic is expected to change in β_k units. This interpretation identifies one weakness of using a linear regression that includes the explanatory variables linearly. It captures only the unconditional mean of the distribution of \mathbf{X} . This can be easily mended by including higher-order polynomials and interactions that would better capture some of the nonlinear relationships across the independent variables and the $\text{RIF}\{y, v(F_Y)\}$. For example, assuming only one exogenous variable, the following specification could be applied:

$$\text{RIF}\{y, v(F_Y)\} = \beta_0 + \beta_1 x + \beta_2 (x - \bar{X})^2 + \varepsilon_i \quad (17)$$

Taking the unconditional expectations, one obtains the following:

$$v(F_Y) = \beta_0 + \beta_1 \bar{X} + \beta_2 E\left\{(X_i - \bar{X})^2\right\} = \beta_0 + \beta_1 \bar{X} + \beta_2 \text{Var}(X)$$

The UPE is now a function of two moments of the unconditional distribution of X , the mean and the variance.¹⁶ An advantage of introducing the term $(x - \bar{X})^2$, as opposed

14. This is a two-step process that is internally done within the community-contributed commands `rifreg` and `xtrifreg`.

15. Firpo, Fortin, and Lemieux (2009) call this change a small location shift in the distribution of \mathbf{X} .

16. While not yet explored, it is possible to use RIFs of other distributional statistics of \mathbf{X} as explanatory variables to better capture how changes in the distribution of \mathbf{X} affect the distribution statistic $v(F_Y)$.

to simply adding x^2 , is that the UPE becomes a function of the mean and variance of X , which one can reasonably assume to change independently from each other.

The estimation of RIF regressions in Stata, under the linearity assumption, is easily implemented using the community-contributed commands `rifreg` (for Gini, variances, and quantiles), `xtrifreg` (for quantiles with one high-dimensional fixed effect), or `rifireg` (for rank-dependent variables). However, there are no commands for the estimation of RIF regressions for other distributional statistics or when two or more high-dimensional fixed effects need to be estimated.

For the estimation of RIF regressions under both scenarios, I introduce the command `rifhdreg`, which is a wrapper that uses the same two-step procedure used in `rifreg`. First, it estimates the corresponding RIF for each observation in the sample of interest for a specific distributional statistic using the previously introduced `rifvar()` function. Second, it uses the RIF as the dependent variable and fits a linear model using the official Stata command `regress` (when no fixed effects are used) or `reghdfe` (Correia 2016) (when fixed effects are used) to fit the RIF-OLS models. The syntax of the command is as follows:

```
rifhdreg depvar [indepvars] [if] [in] [weight], rif(RIF_options)
      [retain(newvar) replace abs(varlist) iseed(str) over(varname)
      rwlogit(varlist) rwprobit(varlist) rwmlogit(varlist) rwmprobit(varlist)
      [ate|att|atu] scale(real) svy regress_options reghdfe_options]
```

`aweights`, `fweights`, `weights`, and `pweights` are allowed; see [U] 11.1.6 **weight**.

The main difference between the `regress` and `reghdfe` commands is that `rifhdreg` requires specifying the distributional statistic of interest with the `rif(RIF_options)` option. See `help rifhdreg` for a description of all options.

`rif(RIF_options)` specifies the statistic of interest, internally estimating the corresponding RIF in a first step. It uses the same syntax presented in table 1. For example, to estimate the RIF regression for the interquantile share ratio, one can type `rifhdreg y x1 x2 x3, rif(iqsr(10 90))`. `rif()` is required.

`retain(newvar)` saves the internally constructed RIF for the restricted sample used in the regression under a new variable name.

`replace`, if used with `retain()`, replaces the internally constructed RIF if the `newvar` specified in `retain()` already exists.

`abs(varlist)` identifies the fixed effects to be absorbed. Each variable listed here represents one set of fixed effects.

When `abs(varlist)` is used, `rifhdreg` calls for `reghdfe` to fit the RIF-OLS model, and all options in `reghdfe` are available. Otherwise, it uses `regress` to fit the model, allowing for all `regress` options.¹⁷

`iseed(str)` is used for the creation of an auxiliary random variable that can be used to obtain replicable estimates for rank-dependent indices when ties exist.

`scale(real)` provides a value for rescaling the dependent variable. It may be useful for statistics like the Lorenz ordinate or quantile shares, which are measures that fall between 0 and 1. In those cases, using `scale(100)` would change the scale to be between 0–100. The default is `scale(1)` (no rescaling).

`rifhdreg` reports OLS asymptotic standard errors by default, but one can request other standard errors allowed in `regress` or `reghdfe` commands, for example, robust and clustered standard errors. Simple survey design standard errors are also available via the option `svy`¹⁸ when no fixed effects are included.

Based on the recommendation provided in Firpo, Fortin, and Lemieux (2009) and the simulations presented in appendix A, bootstrap standard errors should be used when the statistics of interest are the unconditional quantiles or statistics related to unconditional quantiles, or at the very least, robust standard errors should be requested.¹⁹ The reason for this is that RIFs for unconditional quantiles require the estimation of density functions, which are taken as known parameters for the estimation of asymptotic standard errors. For the correct estimation of bootstrap standard errors, one should use the `bootstrap` prefix before the `rifhdreg` command. One can also use the community-contributed command `bs4rw` (Kolenikov 2010) for the estimation of bootstrap standard errors allowing for complex survey design.²⁰

To illustrate the use of this command for the estimation of RIF regressions, table 2 provides the output of RIF regressions for selected distributional statistics, using an excerpt from the Swiss Labor Market Survey available online.²¹ For simplicity, only years of education, years of experience, years of job tenure, sex, and single status are used as explanatory variables. At the bottom of the table, the average RIF is reported as a reference point for the UPE interpretation. Three forms of standard errors are reported: the default or OLS standard errors, robust standard errors, and bootstrap standard errors based on 500 repetitions. The estimation corresponding to the Gini rescales the parameters by 100 to help with the interpretation of the estimated coefficients. All interpretations are given in relative terms with respect to the current levels of inequality to make the results comparable across inequality measures.

17. While all `regress` and `reghdfe` options are permitted, not all of them may be appropriate for the estimation of RIF regressions. I recommend using them with caution.

18. Using this option fits the RIF-OLS model as `svy: regress rif.y x` after the RIF has been estimated using the weights stored as part of the survey design.

19. `rifreg` reports robust standard errors by default.

20. I thank Gopal Trital for pointing this out.

21. The programs used for the replication of the illustrations are available online or by request. The dataset is available online.

Table 2. Determinants of wage inequality

	lnwage iqr(90 10)	wage iqratio(90 10)	wage iqsr(10 90)	wage Gini \times 100	lnwage Variance
Years of education	-0.0154	-0.0542	-0.8580	-0.5715	-0.0357
Default	(0.0127)	(0.0375)	(0.1760)*	(0.2802)+	(0.0082)*
Robust	(0.0154)	(0.0453)	(0.2220)*	(0.3255) \wedge	(0.0117)*
Bootstrap	(0.0188)	(0.0543)	(0.2970)*	(0.3391) \wedge	(0.0114)*
Years of experience	-0.0263	-0.0774	-0.2690	-0.4888	-0.0129
Default	(0.0046)*	(0.0135)*	(0.0636)*	(0.1010)*	(0.0030)*
Robust	(0.0046)*	(0.0135)*	(0.0614)*	(0.1243)*	(0.0026)*
Bootstrap	(0.0052)*	(0.0159)*	(0.0758)*	(0.1267)*	(0.0026)*
Years of job tenure	0.0041	0.0118	0.0490	0.0734	0.0017
Default	(0.0048)	(0.0141)	(0.0662)	(0.1052)	(0.0031)
Robust	(0.0047)	(0.0137)	(0.0533)	(0.1019)	(0.0024)
Bootstrap	(0.0049)	(0.0144)	(0.0510)	(0.0973)	(0.0022)
Is female	0.0438	0.1440	1.4370	3.5025	0.0544
Default	(0.0614)	(0.1810)	(0.8520) \wedge	(1.3534)*	(0.0396)
Robust	(0.0613)	(0.1800)	(0.8280) \wedge	(1.3000)*	(0.0368)
Bootstrap	(0.0671)	(0.1970)	(0.8850)	(1.3531)*	(0.0381)
Age	0.0137	0.0391	-0.0222	0.1806	0.0014
Default	(0.0047)*	(0.0138)*	(0.0651)	(0.1035) \wedge	(0.0030)
Robust	(0.0050)*	(0.0146)*	(0.0819)	(0.1515)	(0.0037)
Bootstrap	(0.0056)+	(0.0165)+	(0.0793)	(0.1466)	(0.0035)
Is single	-0.0637	-0.1810	0.9320	0.0506	0.0439
Default	(0.0709)	(0.2090)	(0.9840)	(1.5627)	(0.0457)
Robust	(0.0617)	(0.1810)	(0.7140)	(1.4090)	(0.0282)
Bootstrap	(0.0652)	(0.1920)	(0.7720)	(1.3658)	(0.0286)
_cons	1.0230	2.9170	18.9000	28.3451	0.7550
Default	(0.1980)*	(0.5830)*	(2.7410)*	(4.3541)*	(0.1270)*
Robust	(0.2280)*	(0.6710)*	(3.5830)*	(4.9969)*	(0.1640)*
Bootstrap	(0.2790)*	(0.8080)*	(4.6740)*	(4.6601)*	(0.1490)*
Avg. RIF	1.0595	2.8844	6.0285	24.6030	0.2818
N	1434	1434	1434	1434	1434

NOTE: Standard errors in parentheses.

 $\wedge p < 0.1$ + $p < 0.05$ * $p < 0.01$

In terms of the standard errors, as expected, OLS standard errors tend to be smaller than robust standard errors for most of the explanatory variables across all models. Bootstrap standard errors, which Firpo, Fortin, and Lemieux (2009) indicate to be the most appropriate for statistical inference, are considerably larger for the regressions that involve quantile statistics. Consistent with the results from the RIF simulations, however, bootstrap errors are similar to the robust standard errors for regressions that focus on the Gini coefficient and variance of log wages.²²

22. While not fully explored, this result suggests that bootstrap standard errors may not be necessary for the analysis of the majority of the distributional statistics. This is in line with Cowell and Flachaire (2007) and Deville (1999), who indicate that RIFs can be used for the estimation of asymptotic standard errors of distributional statistics.

All models have consistent results with different insights for different measures of wage inequality. If the average number of years of education in the population were to increase by one year, the predicted Gini coefficient would drop by just over half a point (2.32% in relative terms), the variance of log wage would drop by 12.7% ($-0.036/0.2818$), and the ratio of wages held by the richest 10% compared with the poorest 10% by 14.2% ($-0.86/6.0285$). The log wage gap and wage ratio between the 90th quantile and 10th quantile may also decline, but the change is not statistically significant. Increasing the number of years of experience seems to reduce wage inequality, whereas aging of the population may have a small effect by increasing the gap between the top and bottom of the wage distribution. This is not observed for other statistics that look across the wage distribution.

Perhaps one of the most challenging variable types to interpret in terms of UPE is categorical variables. On one hand, because RIF and standard RIF regressions are meant to estimate the impact of small changes in the distribution of the independent variables, one should not interpret the coefficients of categorical variables as changes from 0 to 1. Otherwise, the exercise implies a large change in the distribution of the categorical variable, from 0% of observations being classified in that group to 100% being classified in that group, which may introduce a large bias on the predicted UPE.²³ On the other hand, based on (15) and (16), one should analyze the UPE as deviations from the observed unconditional averages. For example, if the proportion of women in the population would increase by 10 percentage points, from 47.6% as currently observed to 57.6%, the ratio of wages held by the richest 10% to the poorest 10% would increase by 2.3% ($1.437/6.0285 \times 0.1$) and the predicted Gini coefficient would increase by 1.4% ($3.502/24.603 \times 0.1$).

To illustrate the capabilities of the `rifhdreg` command to estimate RIF regressions allowing for multiple fixed effects, the model of table 2 is refit, including occupation, age, and single status as fixed effects.²⁴ The results are reported in table 3 using robust and bootstrap standard errors.

23. In the context of RIF regressions, the estimation of the impact of a full change in a binary variable from 0 to 1 is equivalent to what Firpo and Pinto (2016) call inequality treatment effects, which require a slightly different strategy compared with the standard RIF regression approach. Frölich and Melly (2010) implement an estimator of the estimation of quantile treatment effects (`ivqte`) under exogenous and endogenous treatments. Section 4.2 describes advanced options for `rifhdreg` that allow for the estimation of exogenous inequality treatment effects as described in Firpo and Pinto (2016).

24. In this example, the specification could be directly estimated including dummy variables for each age group and occupation, which provides the same results.

Table 3. Determinants of wage inequality

	lnwage iqr(90 10)	wage iqrtio(90 10)	wage iqsr(10 90)	wage Gini	lnwage Variance
Years of education	0.0280	0.0772	0.1490	0.0072	0.0028
Robust	(0.0164)∧	(0.0479)	(0.2110)	(0.0037)+	(0.0120)
Bootstrap	(0.0170)	(0.0509)	(0.2040)	(0.0036)+	(0.0115)
Years of experience	−0.0215	−0.0632	−0.1650	−0.0033	−0.0076
Robust	(0.0047)*	(0.0136)*	(0.0574)*	(0.0012)*	(0.0025)*
Bootstrap	(0.0050)*	(0.0151)*	(0.0623)*	(0.0011)*	(0.0025)*
Years of job tenure	0.0006	0.0010	−0.0440	−0.0005	−0.0023
Robust	(0.0047)	(0.0137)	(0.0574)	(0.0011)	(0.0025)
Bootstrap	(0.0050)	(0.0147)	(0.0566)	(0.0011)	(0.0024)
Is female	0.0790	0.2490	2.0150	0.0355	0.0863
Robust	(0.0627)	(0.1840)	(0.7540)*	(0.0119)*	(0.0356)+
Bootstrap	(0.0698)	(0.2040)	(0.8010)+	(0.0121)*	(0.0352)+
N	1434	1434	1434	1434	1434

NOTES: Standard errors in parentheses.

All models include age, occupation, and single status fixed effects.

∧ $p < 0.1$

+ $p < 0.05$

* $p < 0.01$

In this example, controlling for the age and occupation fixed effects is effectively controlling for all linear and nonlinear relationships these variables have with the dependent variable. In this case, the results suggest that increasing the average number of years of education by one year may increase inequality by about 2.9% regardless of the measure of inequality and holding everything else constant. In contrast, if the number of years of experience increases by one unit, inequality would decline between 1.3% to 2.7%, depending on the inequality measure. Also consistent with the results in table 2, if the proportion of women in the sample increases by 10 percentage points, then inequality measured by the interquantile share ratio, the Gini coefficient, and the log variance will increase by 1.4% to 3.3%.

4.2 Inequality treatment effects: Advanced options

As previously described, standard RIF regressions should not be used to estimate the effect of large changes in the distribution of the independent variables, particularly when considering categorical variables. In fact, Essama-Nssah and Lambert (2012) emphasize that one of the main weaknesses of RIF regressions is that coefficients provide only local approximations of the effect of changes in the distribution of the independent variables.

In an attempt to address this weakness, other studies in the literature, including Rothe (2010), Donald and Hsu (2014), Firpo and Pinto (2016), and Firpo, Fortin, and Lemieux (2018), have proposed methodologies for the estimation of what Firpo and Pinto (2016) call inequality treatment effects. In essence, these methodologies use parametric or nonparametric strategies to obtain inverse probability weights that can be used to identify counterfactual distributions and identify the treatment effects

on distributional statistics. Paraphrasing Firpo and Pinto (2016), the estimation of inequality treatment effects can be described as follows.

Assume there is a joint distribution function $F_{Y_1, Y_0, \mathbf{X}, T}(\cdot)$ that characterizes the relationships between the potential outcomes Y_1 and Y_0 , a set of exogenous explanatory variables \mathbf{X} , and a binary treatment T . If potential outcomes could be observed for everyone in the population, inequality treatment effects could be estimated by simply comparing the distributional statistic of interest v for both potential outcomes as follows:

$$\Delta v = v_1 - v_0 = v(F_{Y_1}) - v(F_{Y_0}) \quad (18)$$

F_{Y_1} and F_{Y_0} represent the unconditional cumulative distributions of the potential outcomes.

Empirically, both potential outcomes are never fully observed. Instead, one is able to observe only the realized outcome depending on whether an individual is in part of the treated group:

$$Y = T \times Y_1 + (1 - T)Y_0$$

According to Firpo and Pinto (2016), under the assumption that the distribution of potential outcomes Y_1 and Y_0 are independent from treatment assignment when conditioning on observed characteristics \mathbf{X} (unconfoundedness) and that there are enough observations so that it is possible to find individuals with similar observed characteristics \mathbf{X} in both groups (overlapping support), treatment effects, as defined in (18), can be identified using the following procedure.

Recall from (9) that the unconditional c.d.f. can also be written by integrating (averaging) the conditional distributions with respect to \mathbf{X} and T :

$$F_{Y_k} = \int F_{Y_k|\mathbf{X}, T} dF_{\mathbf{X}, T} \quad (19)$$

Based on the unconfoundedness assumption, we know that $F_{Y_k|\mathbf{X}, T=1} = F_{Y_k|\mathbf{X}, T=0} = F_{Y_k|\mathbf{X}}$. This implies that (19) can be written solely in terms of the distribution of \mathbf{X} :

$$F_{Y_k} = \int F_{Y_k|\mathbf{X}} dF_{\mathbf{X}}$$

As indicated before, the full distribution of the potential outcomes is never observed. Instead, only the distribution of the realized outcome after treatment has been assigned is available. Using the unconfoundedness assumption again, the observed distribution of the outcome among individuals who are treated and untreated can be defined as

$$F_{Y_k|T=k} = \int_{i \in k} F_{Y_k|\mathbf{X}} dF_{\mathbf{X}|T=k} \quad \text{for } k = 0, 1 \quad (20)$$

which considers only observations that are in the treated or untreated group. Because the distribution implied by (20) is fully observed, it can be used to identify F_{Y_k} using

$$\hat{F}_{Y_k} = \int_{i \in k} F_{Y_k|X} \omega_k(\mathbf{x}) dF_{X|T=k}$$

where $\omega_k(\mathbf{x})$ is a weighting factor such that $\omega_k(\mathbf{X})dF_{X|T=k}$ resembles the unconditional distribution dF_X . This weighting factor can be identified as follows:

$$\begin{aligned} \omega_k(\mathbf{x}) &= \frac{dF_X}{dF_{X|T=k}} = dF_X \times \frac{P(T=k)}{P(T=k|\mathbf{X}=\mathbf{x})dF_X} = \frac{P(T=k)}{P(T=k|\mathbf{X}=\mathbf{x})} \\ \omega_1(\mathbf{x}) &= \frac{P(T=1)}{P(T=1|\mathbf{X}=\mathbf{x})} \quad \text{and} \quad \omega_0(\mathbf{x}) = \frac{1-P(T=1)}{1-P(T=1|\mathbf{X}=\mathbf{x})} \end{aligned}$$

When $k=1$, $P(T=1)$ is the overall probability that an observation is assigned to the treatment group, and $P(T=1|\mathbf{X}=\mathbf{x})$ is the probability that an observation is assigned to the treatment group conditional on the observed characteristics. Following Firpo and Pinto (2016), this probability can be estimated using parametric methods such as probit or logit models, as well as other semiparametric or nonparametric methods.

Once the weights $\omega_1(\mathbf{x})$ and $\omega_0(\mathbf{x})$ have been obtained, the treatment effects defined in (18) can be estimated as follows:

$$\begin{aligned} \Delta \hat{v} &= v(\hat{F}_{Y_1}) - v(\hat{F}_{Y_0}) \\ \Delta \hat{v} &= v \left\{ \int_{i \in 1} F_{y|X} \hat{\omega}_1(\mathbf{x}) dF_{X|T=1} \right\} - v \left\{ \int_{i \in 0} F_{y|X} \hat{\omega}_0(\mathbf{x}) dF_{X|T=0} \right\} \end{aligned}$$

Firpo and Pinto (2016) denominate this as an overall treatment effect, and it is similar to the estimation of average treatment effects.²⁵ While $v(\hat{F}_{Y_1})$ and $v(\hat{F}_{Y_0})$ can be obtained by simply estimating the statistic of interest v using the reweighting factors $\hat{\omega}_1(X)$ and $\hat{\omega}_0(X)$ over the sample of treated and untreated observations, they can also be estimated using RIFs:

$$v(\hat{F}_{Y_k}) = \int_{i \in k} E \left[\text{RIF} \left\{ y, v(\hat{F}_{Y_k}) \right\} | \mathbf{X} = \mathbf{x} \right] \hat{\omega}_k(\mathbf{x}) dF_{X|T=k}(\mathbf{x}) \quad (21)$$

Based on (21), treatment effects can be estimated using RIF regressions, fitting the model

$$T \times \text{RIF} \left\{ y, v(\hat{F}_{Y_1}) \right\} + (1-T) \times \text{RIF} \left\{ y, v(\hat{F}_{Y_0}) \right\} = b_0 + b_1 T + \varepsilon \quad (22)$$

25. Following Firpo and Pinto (2016), one can also estimate average treatment effects on the treated or average treatment effects on the untreated by simply modifying the weight functions. For average treatment effects on the treated, $\omega_1(\mathbf{X}) = 1$ and $\omega_0(\mathbf{X}) = \{1 - P(T=1)\} / \{P(T=1)\} \{P(T=1|\mathbf{X})\} / \{1 - P(T=1|\mathbf{X})\}$. For the average treatment effects on the untreated, $\omega_1(\mathbf{X}) = \{P(T=1)\} / \{1 - P(T=1)\} \{1 - P(T=1|\mathbf{X})\} / \{P(T=1|\mathbf{X})\}$ and $\omega_0(\mathbf{X}) = 1$.

using weighted least squares with weights equal to $\widehat{\omega}(\mathbf{x}) = T\widehat{\omega}_1(\mathbf{x}) + (1 - T)\widehat{\omega}_0(\mathbf{x})$, and where the dependent variable is constructed by combining the RIFs based on the reweighted cumulative distribution functions \widehat{F}_{Y_1} and \widehat{F}_{Y_0} , which are estimated separately for observations in the treated and untreated groups.

Applying conditional expectations with respect to T on both sides of (22) and taking the derivative with respect to T , it is easy to see that the overall inequality treatment effect (average treatment effect) is equal to b_1 .²⁶

$$\begin{aligned} T \times E \left[\text{RIF} \left\{ y, v \left(\widehat{F}_{Y_1} \right) \right\} | T = 1 \right] &+ (1 - T) \\ &\times E \left[\text{RIF} \left\{ y, v \left(\widehat{F}_{Y_0} \right) \right\} | T = 0 \right] = b_0 + b_1 T \\ T \times v \left(\widehat{F}_{Y_1} \right) &+ (1 - T) \times v \left(\widehat{F}_{Y_0} \right) = b_0 + b_1 T \\ v \left(\widehat{F}_{Y_1} \right) - v \left(\widehat{F}_{Y_0} \right) &= b_1 \end{aligned}$$

The advantage of using RIF regressions for the estimation of treatment effects using specifications similar to (22) is that one can also control for differences in the distribution of characteristics directly by including them as control variables in the model specification. In addition, if the variable T has more than two categories, (22) can still be estimated by using RIFs constructed using all categories in T , and coefficients associated with T can be interpreted as treatment effects, while coefficients of other variables would be interpreted as UPE that are averaged across the distributions defined by T .²⁷

The command `rifhdreg` offers three advanced options that can be used to estimate inequality treatment effects using models like the one just described.

The first option is `over(varname)`. It is used to specify a variable over which the RIF will be estimated, just as the option `by()` is used in the `rifvar()` function. For example, if one uses the command `rifhdreg y x i.T, over(T) rif(q(10))`, where T is a categorical variable, the command will first estimate the 10th quantile RIF for y for all groups defined by T and use that as the dependent variable in the regression model. When the variable used in `over(varname)` is binomial, the regression can also be seen as the OLS alternative to OB decomposition.

`rwlogit(varlist)` and `rwprobit(varlist)` are options used to specify the estimation of the reweighting factors $\omega_k(\mathbf{x})$ using a logit or probit model and using the variable specified in `over()` as the dependent variable. These options can be used to estimate distributional treatment effects under the exogeneity assumption, using weighted least

26. This simple framework can also be extended for the estimation of overall multivalued treatment effects as described in Cattaneo (2010) and Cattaneo, Drukker, and Holland (2013), based on an inverse probability weighting strategy. Thank you to David Drukker for suggesting this extension to the command.

27. Similarly to Cattaneo, Drukker, and Holland (2013), a multinomial logit or probit can be used to estimate the appropriate generalized propensity scores, which can in turn be used for the estimation of multivalued treatment effects. Also see references on `teffects ipw` and `teffects multivalued` for estimation of treatment effects using the official Stata command.

squares based on (22), as described in Firpo and Pinto (2016). When either of these options is used, the user should also specify the type of treatment effect that will be estimated. The options are average treatment effect, `ate` (default); treatment effect on the treated, `att`; or treatment effect on the untreated, `atu`. When using this option, a new variable will be created under the name `_wipw_` containing the estimated $\omega(\mathbf{X})$. Using `rwlogit()` or `rwprobit()` in combination with the inclusion of independent variables in the model specification can be compared with the inverse-probability weighted regression adjustment estimator of treatment effects.

`rwmllogit(varlist)` and `rwmlprobit(varlist)` are options used to specify the estimation of the reweighting factors $\omega_k(\mathbf{x})$ using a multinomial logit or multinomial probit model, using the variable specified in `over()` as the dependent variable. These options can be used to estimate distributional multivalued treatment effects under the exogeneity assumption using weighted least squares. When using any of these options, only the average treatment effect (`ate`) can be estimated.

To provide an example showing the estimation of inequality treatment effects using `rifhdreg`'s advanced options, table 4 (below) provides the estimated overall treatment effects following the same specification of the models presented in table 2, but where the corresponding RIFs are calculated over the treatment variable being a woman. Four specifications are used: i) no controls, ii) no controls but using the reweighting adjustment (IPW), iii) adding controls as in table 2, but without the reweighting adjustment, and iv) adding both controls and the reweighting adjustment. The conditional probabilities of treatment are estimated with a probit model that uses the same variables used in the outcome model.

The first column in table 4 provides the estimation of overall treatment effects without using any controls. These results provide an estimate of the raw distributional gap between women and men and are most likely to provide biased estimates for treatment effects. Overall, this result indicates that women earn on average about 17.4% lower wages, with a larger gap (24.8%) at the 10th quantile but a somewhat smaller gap (13.9%) at the 90th quantile. This column also suggests that there is somewhat higher wage inequality among women.

The second, third, and fourth columns, which control for differences in characteristics in various ways, suggest that the negative treatment effects of gender at the means and the 10th and 90th quantiles are smaller than the raw gaps would suggest, albeit only marginally. In terms of inequality, after controlling for other characteristics, the results suggest that the treatment effect of gender may not increase inequality as much as the raw gender gap suggests. The three methods used to control for differences in characteristics provide similar results, with the largest differences observed for the treatment effect on the Gini coefficient and the interquantile ratio.

Table 4. Estimation of overall treatment effects of gender

Treatment effect: Female	No controls	No controls +IPW	Controls	Controls +IPW
Mean (<code>lnwage</code>)	-0.1735	-0.1400	-0.1290	-0.1376
Robust	(0.0279)*	(0.0290)*	(0.0247)*	(0.0251)*
Bootstrap	(0.0271)*	(0.0243)*	(0.0238)*	(0.0245)*
10th quantile (<code>lnwage</code>)	-0.2486	-0.2220	-0.1977	-0.2191
Robust	(0.0448)*	(0.0489)*	(0.0441)*	(0.0463)*
Bootstrap	(0.0487)*	(0.0501)*	(0.0484)*	(0.0499)*
90th quantile (<code>lnwage</code>)	-0.1386	-0.1233	-0.1270	-0.1219
Robust	(0.0376)*	(0.0413)*	(0.0391)*	(0.0394)*
Bootstrap	(0.0503)*	(0.0546)+	(0.0512)+	(0.0544)+
<code>iqr(90 10) (lnwage)</code>	0.1100	0.0987	0.0710	0.0973
Robust	(0.0544)+	(0.0596)^	(0.0572)	(0.0588)^
Bootstrap	(0.0697)	(0.0767)	(0.0725)	(0.0764)
<code>iqratio(90 10) (wage)</code>	0.3028	0.2733	0.1927	0.2690
Robust	(0.1567)^	(0.1747)	(0.1646)	(0.1723)
Bootstrap	(0.2024)	(0.2198)	(0.2087)	(0.2189)
<code>iqsr(90 10) (wage)</code>	2.3540	1.5549	1.4275	1.5335
Robust	(0.9301)+	(0.8401)^	(0.8904)	(0.8141)^
Bootstrap	(0.9602)+	(0.8434)^	(0.8723)	(0.8481)^
Gini (<code>wage</code>)	4.6632	3.7194	3.4512	3.7124
Robust	(1.3548)*	(1.3290)*	(1.3331)*	(1.3204)*
Bootstrap	(1.4118)*	(1.3778)*	(1.3810)+	(1.3845)*
Variance (<code>lnwage</code>)	0.0957	0.0599	0.0540	0.0590
Robust	(0.0381)+	(0.0370)	(0.0366)	(0.0361)
Bootstrap	(0.0392)+	(0.0376)	(0.0383)	(0.0378)
<i>N</i>	1434	1434	1434	1434

NOTE: Robust and bootstrap standard errors, based on 500 repetitions, in parentheses. All models provide estimations for overall treatment effects of gender. Four specifications are used, combining controls and reweighting adjustment.

^ $p < 0.1$

+ $p < 0.05$

* $p < 0.01$

5 RIF decomposition: `oaxaca_rif`

As previously described, one of the main advantages of standard RIF regressions described in section 4.1 is that they can be easily used to analyze how small changes in the distribution of independent characteristics will affect the distributional statistic v . Furthermore, the modified RIF regression described in section 4.2 can also be used to estimate inequality treatment effects caused by an exogenous treatment.

However, there is often interest in analyzing what factors explain the differences in the distribution between two groups. The OB decomposition is one of the most extensively used methodologies in labor economics that aims to analyze outcome differences between two groups (Blinder 1973; Oaxaca 1973). These differences are characterized as functions of differences in characteristics (composition effect) and differences in coefficients associated with those characteristics (wage structure effect).

While the original methodology was created to analyze differences of outcome means, several articles that followed provided extensions and refinements to extend the analysis to other distributional statistics (see Fortin, Lemieux, and Firpo [2011] for a review). Additionally, under the assumptions of conditional independence (unconfoundedness) and overlapping support, the aggregate structure effect can be identified and interpreted as a treatment effect.²⁸

Firpo, Fortin, and Lemieux (2018) describe the use of RIF regressions, in combination with a reweighted strategy (DiNardo, Fortin, and Lemieux 1996), as a feasible methodology for decomposing differences in distributional statistics beyond the mean. This is referred to as RIF decomposition. This methodology has three advantages compared with other strategies in the literature: the simplicity of its implementation, the possibility of obtaining detailed contributions of individual covariates on the aggregate decomposition,²⁹ and the possibility of expanding the analysis to any statistic for which a RIF can be defined. This strategy can be described as follows.

Assume there is a joint distribution function that describes all relationships between the dependent variable Y , the exogenous characteristics \mathbf{X} , and the categorical variable T that identifies the group membership of the population ($f_{Y,X,T}$). The cumulative distribution of Y conditional on T can be written as

$$F_{Y|T=k} = \int F_{Y|X,T=k} dF_{X|T=k}$$

To analyze the differences between groups 0 and 1, the cumulative conditional distribution of Y can be used to calculate the gap in the distributional statistic v :

$$\begin{aligned} \Delta v &= v_1 - v_0 = v(F_{Y|T=1}) - v(F_{Y|T=0}) \\ \Delta v &= v \left(\int F_{Y|X,T=1} dF_{X|T=1} \right) - v \left(\int F_{Y|X,T=0} dF_{X|T=0} \right) \end{aligned} \quad (23)$$

28. These are the same assumptions used by Firpo and Pinto (2016) for the identification of the treatment effects.

29. Akin to the standard regression analysis, the identification of the detailed contribution of covariates also requires the zero-conditional mean assumption. In other words, any other variable not accounted for in the model has a distribution that is independent from the measured characteristics \mathbf{X} .

From (23), it is easy to see that differences in the statistics of interest Δv will arise because of differences in the distribution of \mathbf{X} ($dF_{X|T=1} \neq dF_{X|T=0}$) or because of differences in the relationships between Y and \mathbf{X} ($dF_{Y|X,T=1} \neq dF_{Y|X,T=0}$). In the context of the standard OB decomposition, this is equivalent to comparing differences in average characteristics and differences in coefficients.

To identify how important differences in characteristics (composition effect) and differences in coefficients (wage structure effect) are for explaining the overall gap in the distributional statistic v , it is necessary to create a counterfactual scenario. Define the counterfactual statistic v_c as follows:

$$v_c = v(F_Y^c) = v\left(\int F_{Y|X,T=0} dF_{X|T=1}\right)$$

Using this counterfactual, the gap in the distribution statistic v can be disaggregated into two components:

$$\Delta v = \underbrace{v_1 - v_c}_{\Delta v_S} + \underbrace{v_c - v_0}_{\Delta v_X}$$

Δv_X reflects the gap attributed to differences in characteristics, and Δv_S reflects the differences attributed to the relationships between Y and \mathbf{X} (structure effect). The difficulty of this strategy lies in the identification of the counterfactual statistic v_c because the combination of characteristics and outcomes is not observed in the data. Based on the review in Fortin, Lemieux, and Firpo (2011), two broad strategies have been suggested for the identification of the counterfactual statistic v_c . The first strategy follows the standard OB decomposition, using linear regressions and their approximations to identify v_c . Specifically, following (15), separate RIF regressions can be estimated for each group, so the counterfactual statistic can be identified as follows:

$$\begin{aligned} v_1 &= E[\text{RIF}\{y, v(F_{Y|T=1})\}] = \bar{\mathbf{X}}^1{}' \hat{\boldsymbol{\beta}}^1 \\ v_0 &= E[\text{RIF}\{y, v(F_{Y|T=0})\}] = \bar{\mathbf{X}}^0{}' \hat{\boldsymbol{\beta}}^0 \\ v_c &= \bar{\mathbf{X}}^1{}' \hat{\boldsymbol{\beta}}^0 \end{aligned} \tag{24}$$

This alternative mirrors the standard OB decomposition, where $\Delta v_X = (\bar{\mathbf{X}}^1 - \bar{\mathbf{X}}^0)' \boldsymbol{\beta}^0$ and $\Delta v_S = \bar{\mathbf{X}}^1{}' (\hat{\boldsymbol{\beta}}^1 - \hat{\boldsymbol{\beta}}^0)$. The main disadvantage of this strategy, discussed in Barsky et al. (2002) in the context of conditional means, is that the counterfactual statistic v_c may be incorrectly identified if the model is misspecified³⁰ or if the local approximation obtained using RIF cannot be extended beyond local extrapolations. The alternative is to use a semiparametric reweighting approximation, as discussed in Barsky et al. (2002) and DiNardo, Fortin, and Lemieux (1996), to identify the counterfactual distribution $F_{Y|X,T=0} dF_{X|T=1}$ based on the observed data. This procedure can be described as follows.

30. The concept of misspecification here also includes the idea of accounting for changes in the whole distribution of \mathbf{X} , not only the mean.

The problem of identifying the counterfactual scenario is that the distribution of outcomes and characteristics that the counterfactual distribution $F_{Y|X}^C$ implies cannot be directly observed. However, from an abstract point of view, it is possible to obtain an approximation for the counterfactual distribution by multiplying the observed distribution of characteristics $dF_{X|T=0}$ with a factor $\omega(\mathbf{X})$, so it resembles the distribution $dF_{X|T=1}$:

$$F_Y^C = \int F_{Y|X,T=0} dF_{X|T=1} \cong \int F_{Y|X,T=0} dF_{X|T=0} \omega(\mathbf{X})$$

Using the Bayes rule, the reweighting factor $\omega(\mathbf{X})$ can be identified as follows:

$$\begin{aligned} \omega(\mathbf{X}) &= \frac{dF_{X|T=1}}{dF_{X|T=0}} = \frac{dF_{T=1|X} dF_X}{dF_{T=1}} \frac{dF_{T=0}}{dF_{T=0|X} dF_X} = \frac{dF_{T=0}}{dF_{T=1}} \frac{dF_{T=1|X}}{dF_{T=0|X}} \\ &= \frac{1-P}{P} \frac{P(T=1|\mathbf{X})}{1-P(T=1|\mathbf{X})} \end{aligned}$$

p is the proportion of people in group $T=1$, and $P(T=1|\mathbf{X})$ is the conditional probability of someone with characteristics \mathbf{X} being part of group 1. In other words, to identify the counterfactual distribution $F_{Y|X}^C$, one can estimate the reweighting factor $\omega(\mathbf{X})$ using parametric or nonparametric methods to estimate the conditional probability $P(T=1|\mathbf{X})$. As described in Firpo and Pinto (2016), in practice, a probit or logit model can be used to estimate this conditional probability.

Once these reweighting factors are obtained, (24) is estimated using weighted least squares:

$$v_c = E [\text{RIF} \{y, v(F_Y^C)\}] = \bar{\mathbf{X}}^{c'} \hat{\beta}^c \quad (25)$$

The decomposition components are now defined as

$$\Delta v = \underbrace{\bar{\mathbf{X}}^{1'} (\hat{\beta}_1 - \hat{\beta}_c)}_{\Delta v_s^p} + \underbrace{(\bar{\mathbf{X}}^1 - \bar{\mathbf{X}}^c)' \hat{\beta}_c}_{\Delta v_s^e} + \underbrace{(\bar{\mathbf{X}}^c - \bar{\mathbf{X}}^0)' \hat{\beta}_0}_{\Delta v_X^p} + \underbrace{\bar{\mathbf{X}}^{c'} (\hat{\beta}_c - \hat{\beta}_0)}_{\Delta v_X^e} \quad (26)$$

The components $\Delta v_s^p + \Delta v_s^e$ correspond to the OB aggregate wage structure effect, whereas $\Delta v_X^p + \Delta v_X^e$ correspond to the aggregate composition effect. These two components are further decomposed into a pure wage structure (Δv_s^p) and pure composition effect (Δv_X^p), plus two components that can be used to assess the overall fitness of the model. Δv_s^e is the reweighting error that is used to evaluate the quality of the reweighting strategy and is expected to go to 0 in large samples. A large significant reweighting error implies that the counterfactual is not well identified and that the specification of the probit or logit model used for the estimation of reweighting factors may need to be modified.

Δv_X^e is the specification error and is used to assess the quality of the model specification and the RIF approximation. A large and significant specification error may be an indication that the RIF regression is misspecified or that the RIF is providing a poor approximation to the distributional statistic v .

Implementing OB decomposition in Stata is simple. The most popular command used for this type of analysis is the community-contributed command `oaxaca` (Jann 2008), which can be used for many of the extensions that have been developed for analyzing average differences across groups. Extending the OB decomposition analysis to statistics other than the mean can easily be done by carefully calculating RIFs for the conditional distributions and using them as the dependent variable with the `oaxaca` command. However, no formal implementation of the estimation of RIF decompositions and the hybrid reweighted RIF decomposition is currently available.

For estimating these two types of decompositions, I present the `oaxaca_rif` command, which is a wrapper around `oaxaca` that uses the processes suggested in Firpo, Fortin, and Lemieux (2018) for the estimation of the standard RIF decomposition [following (24)] or the hybrid reweighted decomposition (25).

The syntax of the command is as follows:

```
oaxaca_rif depvar [indepvars] [if] [in] [weight], by(groupvar)
rif(RIF_options) [relax noisily swap wgt(#) scale(real)
cluster(varname) retain(newvar) replace s2var(varlist) rwlogit(varlist)
rwprobit(varlist) iseed(str)]
```

`aweights`, `fweights`, `iwweights`, and `pweights` are allowed; see [U] 11.1.6 **weight**.

Parallel to the `rifhdreg` command, `oaxaca_rif` requires the `rif(RIF_options)` option to define the distributional statistic to be used for the decomposition analysis. Internally, it calls on `rifvar()` to estimate the RIF for each group defined by `by()` and follows (24) or (25) to identify the counterfactuals and implement the decomposition. `by()` is required.

The internal syntax of `oaxaca_rif` allows the use of many features available in `oaxaca` that are appropriate for the estimation of RIF decompositions. These options include `relax`, when some variables show no variation for one of the outcome models, and `noisily`, which displays the intermediate outcome models. In contrast with the `oaxaca` command, when `noisily` is used the coefficients of the outcome models are stored as matrices in `e()`. The option `normalize()` and aggregation of subset of variables using the syntax `[name:] varlist` are also possible.

`swap` requests the gap to be estimated in the opposite order (`group2 – group1`). The default is to estimate the distributional statistic gap between observations with the lowest value in the grouping variable minus the observations in the group with the highest value (`group1 – group2`).

`wgt(#)` defines the counterfactual distribution. The default is `wgt(0)`, which identifies the decomposition according to (24). Using `wgt(1)` instead uses $v_c = \bar{\mathbf{X}}^{0'} \hat{\beta}^1$ as the counterfactual. Values other than 0 or 1 are not allowed.

`scale(real)` provides a value for rescaling the dependent variable. It may be useful for statistics like the Lorenz ordinate or quantile shares, which are measures that fall between 0 and 1. The default is `scale(1)` (no rescaling).

For standard errors, the default is to report robust standard errors, equivalent to using the `robust` option in the `oaxaca` command.³¹ When the reweighted decomposition is requested, the command reports robust standard errors clustered at the individual level. This is done because for one of the internally estimated decompositions [see (25)] the same sample is used for both groups. The option `cluster(varname)` supersedes the individual cluster option. Weights are allowed when robust or clustered standard errors are used.

`retain(newvar)` stores the internally generated RIF into a new variable or replaces the existing variable if `replace` is used. For the reweighted decomposition, `retain()` does not generate the RIF for the counterfactual option.

Similarly to the limitations of RIF regressions, the inclusion of linear terms in the model specification accounts only for differences in average characteristics. As shown in (17), one can control for differences in the dispersion of characteristics by simply adding $(x - \bar{X})^2$ to the model specification. Alternatively, one can use the option `s2var(varlist)` to request centered quadratic terms to be calculated for all listed variables and to be added to the outcome model. This may reduce problems of error due to model misspecification by controlling for other aspects of the distribution of the independent variables. This is the preferred method when estimating reweighted RIF decomposition.³²

`rwlogit(varlist)` and `rwprobit(varlist)` are options used to specify the estimation of the reweighting factors using a logit or probit model. When one of these options is used, `oaxaca_rif` estimates the reweighted RIF decomposition.

For the reweighted standard decomposition, `oaxaca_rif` first fits the probability model, then estimates the reweighting factor $\omega(\mathbf{X})$, and then obtains the RIFs for the three scenarios. The decomposition output is obtained from two separate decompositions to identify the four components detailed in (26).

The variables included in the option `rwlogit(varlist)` or `rwprobit(varlist)` may or may not be the same as the ones used in the specification of the main model. While factor notation cannot be used as part of the outcome model specification, it is permitted for the specification of the probability models. Adding higher-order polynomials and interactions for the estimation of the conditional probability will improve the quality of the reweighting strategy, thus improving the balance in the counterfactual scenario, but may create problems of overfitting and violation of the overlapping assumption.

As shown for the case of RIF regressions, asymptotic and robust standard errors may not be appropriate for the decomposition of statistics related to quantiles or the Atkinson index. Furthermore, because the reweighting factors $\omega(\mathbf{X})$ are estimated vari-

31. Robust standard errors for the `oaxaca` command are not available in all versions of the command. For `oaxaca_rif` to work properly, be sure to have the latest version, which at the time of writing this article is version 4.0.5.

32. I thank an anonymous referee for suggesting the inclusion of this option to the command.

ables, standard errors of the decomposition components need to be adjusted. Given the complexity of estimating asymptotic standard errors in the framework of RIF decompositions, the suggested alternative is to use bootstrap standard errors throughout, which can be obtained using the `bootstrap` prefix. Bootstrap standard errors cannot be used in combination with weights, but one can use the community-contributed command `bs4rw` to account for weights and survey design.

For illustration, I present in table 5 a simple exercise decomposing the same five measures of wage inequality used for the RIF regression example, analyzing differences by gender. Only results for the reweighted RIF decomposition are shown. A logit probability model is used for estimating the reweighting factors, using the same specification as the outcome model.

Table 5. Determinants of gender wage inequality

	lnwage iqr(90 10)	wage iqratio(90 10)	wage iqsr(10 90)	wage Gini \times 100	lnwage Variance
Overall:					
Men	0.9520	2.5900	4.8230	22.0724	0.2290
Robust	(0.0331)*	(0.0861)*	(0.3980)*	(0.6764)*	(0.0216)*
Bootstrap	(0.0336)*	(0.0882)*	(0.2290)*	(0.6545)*	(0.0204)*
Counterfactual:					
Women Xs with Men Bs	1.0830	2.9540	6.2010	25.8587	0.2850
Robust	(0.0473)*	(0.1520)*	(0.5980)*	(1.1785)*	(0.0269)*
Bootstrap	(0.0575)*	(0.1690)*	(0.6500)*	(1.1754)*	(0.0267)*
Women	1.0620	2.8930	7.1770	26.7356	0.3240
Robust	(0.0432)*	(0.1310)*	(0.8400)*	(1.1733)*	(0.0314)*
Bootstrap	(0.0601)*	(0.1800)*	(0.9900)*	(1.2255)*	(0.0326)*
Total difference	-0.1100	-0.3030	-2.3540	-4.6632	-0.0957
Robust	(0.0544)+	(0.1570) \wedge	(0.9300)+	(1.3544)*	(0.0381)+
Bootstrap	(0.0697)	(0.2020)	(1.1020)	(1.4118)	(0.0392)
Total composition effect	0.0210	0.0610	-0.9770	-0.9020	-0.0391
Robust	(0.0274)	(0.0876)	(0.3170)*	(0.3795)+	(0.0087)*
Bootstrap	(0.0374)	(0.1120)	(0.3910)+	(0.3889)	(0.0118)
Total wage structure	-0.1310	-0.3640	-1.3770	-3.7863	-0.0567
Robust	(0.0570)+	(0.1730)+	(0.7100)+	(1.3472)*	(0.0340) \wedge
Bootstrap	(0.0686) \wedge	(0.1960) \wedge	(0.7450) \wedge	(1.3416)*	(0.0333) \wedge
Total composition effect:					
Specification error	0.0276	0.0777	0.4130	-0.0252	0.0062
Robust	(0.0226)	(0.0721)	(0.1930)+	(0.1688)	(0.0068)
Bootstrap	(0.0310)	(0.0913)	(0.1830)+	(0.1402)	(0.0049)
Pure composition effect	-0.0065	-0.0166	-1.3900	-0.9020	-0.0453
Robust	(0.0194)	(0.0589)	(0.3600)*	(0.3795)+	(0.0133)*
Bootstrap	(0.0326)	(0.0952)	(0.5610)+	(0.4331)+	(0.0154)*
Years of education	0.0014	0.0067	-0.7820	-0.4892	-0.0263
Robust	(0.0131)	(0.0399)	(0.2530)*	(0.2832) \wedge	(0.0106)+
Bootstrap	(0.0198)	(0.0584)	(0.2860)*	(0.3388)	(0.0127)+
Years of experience	-0.0095	-0.0277	-0.5810	-0.4267	-0.0187
Robust	(0.0078)	(0.0238)	(0.1650)*	(0.1731)+	(0.0061)*
Bootstrap	(0.0141)	(0.0417)	(0.2670)+	(0.2397) \wedge	(0.0094)+
Years of job tenure	0.0016	0.0043	-0.0271	0.0139	-0.0003
Robust	(0.0143)	(0.0433)	(0.1970)	(0.2951)	(0.0078)
Bootstrap	(0.0175)	(0.0532)	(0.2130)	(0.3343)	(0.0082)

Continued on next page

	lnwage iqr(90 10)	wage iqratio(90 10)	wage iqsr(10 90)	wage Gini \times 100	lnwage Variance
Total wage structure:					
Reweighting error	0.0004	0.0013	-0.1180	-0.1192	-0.0052
Robust	(0.0047)	(0.0151)	(0.1600)	(0.1792)	(0.0072)
Bootstrap	(0.0046)	(0.0147)	(0.0704) \wedge	(0.1053)	(0.0035)
Pure wage structure	-0.1310	-0.3650	-1.2590	-3.6672	-0.0515
Robust	(0.0582)+	(0.1770)+	(0.6850) \wedge	(1.3374)*	(0.0328)
Bootstrap	(0.0691) \wedge	(0.1980) \wedge	(0.7070) \wedge	(1.3527)*	(0.0334)
Years of education	-0.1320	-0.6040	-0.3510	6.8727	0.0049
Robust	(0.3400)	(1.0360)	(3.9270)	(7.1192)	(0.2180)
Bootstrap	(0.3580)	(1.0850)	(3.9820)	(6.8546)	(0.1940)
Years of experience	-0.3400	-0.9160	-0.4680	-3.4572	-0.0467
Robust	(0.1010)*	(0.3130)*	(1.1230)	(2.0696) \wedge	(0.0555)
Bootstrap	(0.1090)*	(0.3340)*	(1.1390)	(2.0156) \wedge	(0.0525)
Years of job tenure	0.1240	0.3150	0.4840	2.3062	0.0256
Robust	(0.0823)	(0.2550)	(0.6800)	(1.5275)	(0.0293)
Bootstrap	(0.0865)	(0.2680)	(0.7180)	(1.5876)	(0.0299)
_cons	0.2170	0.8410	-0.9240	-9.3890	-0.0352
Robust	(0.3780)	(1.1510)	(4.8110)	(7.9899)	(0.2640)
Bootstrap	(0.4110)	(1.2260)	(6.1310)	(7.9680)	(0.2320)
N	1434	1434	1434	1434	1434
Men	751	751	751	751	751
Women	683	683	683	683	683

NOTES: Standard errors in parentheses. Bootstrap standard errors estimated using 500 repetitions.

$\wedge p < 0.1$

+ $p < 0.05$

* $p < 0.01$

In general, all models suggest that inequality is higher among women compared with men. Inequality among women is between 11.5% to almost 48% higher than for men. The smallest gaps are measured for the interquartile ratio, while the largest is observed for the interquantile share ratio of wages. Controlling for differences in the distribution of years of education, tenure, and experience slightly reduces the wage inequality by almost half when looking through the wage distribution (interquantile share ratio and variance of log wages) but seems to have a marginal but increasing effect on the interquantile range and interquantile ratio.

The specification error is small but statistically significant for the interquantile share ratio model, suggesting that higher polynomials should be included to reduce the specification bias. The reweighting error suggests a good-quality reweighting because it is small and nonstatistically significant across models with the exception of the interquantile share ratio, where differences in years of experience are still important for the model.

The detailed decomposition effect suggests that women experience larger wage inequality because they have lower levels of education, experience, and job tenure (improvements in those characteristics seem to reduce wage inequality). The wage structure effect also has a negative contribution to explaining inequality because the inequality-reducing effects of education, experience, and job tenure are smaller for women compared with men.

Similarly to the results for RIF regressions, robust standard errors tend to be smaller than bootstrap standard errors, especially for statistics related to quantiles (in table 4, columns 1–3). Nevertheless, because of the two-step nature of the reweighting process, bootstrap standard errors are still recommended.

6 Conclusions

IFs and RIFs are important statistical tools that can be used to analyze the robustness of statistics to outliers and obtain asymptotic standard errors of otherwise complex distributional statistics (Cowell and Flachaire 2007; Deville 1999). Firpo, Fortin, and Lemieux (2009) expand on this literature proposing the use of RIFs in the context of regression and decomposition analysis. This is a simple strategy to analyze unconditional partial effects on any distributional statistics for which a RIF can be obtained.

This article revises the intuition behind the IF and RIF and briefly discusses the setup under which they can be used for regression and decomposition analysis. To facilitate the implementation of these strategies and make RIFs an easy-to-use tool for the applied econometrician, I introduce one function and two new commands: `rifvar()`, for the estimation of RIFs for a large set of distributional statistics; `rifhdreg`, for the estimation of RIF regressions and inequality treatment effects; and `oaxaca_rif`, for the estimation of standard and reweighted OB decompositions for statistics beyond the mean.

7 Programs and supplemental materials

To install a snapshot of the corresponding software files as they existed at the time of publication of this article, type

```
. net sj 20-1
. net install st0588      (to install program files, if available)
. net get st0588          (to install ancillary files, if available)
```

8 References

- Barsky, R., J. Bound, K. K. Charles, and J. P. Lupton. 2002. Accounting for the black–white wealth gap: A nonparametric approach. *Journal of the American Statistical Association* 97: 663–673. <https://doi.org/10.1198/016214502388618401>.
- Blinder, A. S. 1973. Wage discrimination: Reduced form and structural estimates. *Journal of Human Resources* 8: 436–455. <https://doi.org/10.2307/144855>.
- Borgen, N. T. 2016. Fixed effects in unconditional quantile regression. *Stata Journal* 16: 403–415. <https://doi.org/10.1177/1536867X1601600208>.
- Cattaneo, M. D. 2010. Efficient semiparametric estimation of multi-valued treatment effects under ignorability. *Journal of Econometrics* 155: 138–154. <https://doi.org/10.1016/j.jeconom.2009.09.023>.

- Cattaneo, M. D., D. M. Drukker, and A. D. Holland. 2013. Estimation of multivalued treatment effects under conditional independence. *Stata Journal* 13: 407–450. <https://doi.org/10.1177/1536867X1301300301>.
- Chung, C., and P. Van Kerm. 2018. Foreign workers and the wage distribution: What does the influence function reveal? *Econometrics* 6: 41. <https://doi.org/10.3390/econometrics6030041>.
- Correia, S. 2016. A feasible estimator for linear models with multi-way fixed effects. <http://scoreia.com/research/hdfe.pdf>.
- Cowell, F. A., and E. Flachaire. 2007. Income distribution and inequality measurement: The problem of extreme values. *Journal of Econometrics* 141: 1044–1072. <https://doi.org/10.1016/j.jeconom.2007.01.001>.
- . 2015. Statistical methods for distributional analysis. In *Handbook of Income Distribution*, vol. 2, ed. A. B. Atkinson and F. Bourguignon, 359–465. The Netherlands: Elsevier.
- Davies, J. B., N. M. Fortin, and T. Lemieux. 2017. Wealth inequality: Theory, measurement and decomposition. *Canadian Journal of Economics* 50: 1224–1261. <https://doi.org/10.1111/caje.12313>.
- Déville, J.-C. 1999. Variance estimation for complex statistics and estimators: Linearization and residual techniques. *Survey Methodology* 25: 193–203.
- DiNardo, J., N. M. Fortin, and T. Lemieux. 1996. Labor market institutions and the distribution of wages, 1973–1992: A semiparametric approach. *Econometrica* 64: 1001–1044. <https://doi.org/10.2307/2171954>.
- Donald, S. G., and Y.-C. Hsu. 2014. Estimation and inference for distribution functions and quantile functions in treatment effect models. *Journal of Econometrics* 178: 383–397. <https://doi.org/10.1016/j.jeconom.2013.03.010>.
- Efron, B. 1982. *The Jackknife, the Bootstrap and Other Resampling Plans*. Philadelphia: Society for Industrial and Applied Mathematics.
- Essama-Nssah, B., and P. J. Lambert. 2012. Influence functions for policy impact analysis. In *Inequality, Mobility and Segregation: Essays in Honor of Jacques Silber*, ed. J. A. Bishop and R. Salas, 135–159. Bingley, UK: Emerald.
- Firpo, S., N. M. Fortin, and T. Lemieux. 2009. Unconditional quantile regressions. *Econometrica* 77: 953–973. <https://doi.org/10.3982/ECTA6822>.
- Firpo, S., and C. Pinto. 2016. Identification and estimation of distributional impacts of interventions using changes in inequality measures. *Journal of Applied Econometrics* 31: 457–486. <https://doi.org/10.1002/jae.2448>.
- Firpo, S. P., N. M. Fortin, and T. Lemieux. 2018. Decomposing wage distributions using recentered influence function regressions. *Econometrics* 6: 28. <https://doi.org/10.3390/econometrics6020028>.

- Fortin, N., T. Lemieux, and S. Firpo. 2011. Decomposition methods in economics. In *Handbook of Labor Economics*, vol. 4A, ed. O. Ashenfelter and D. Card, 1–102. Amsterdam: Elsevier.
- Frölich, M., and B. Melly. 2010. Estimation of quantile treatment effects with Stata. *Stata Journal* 10: 423–457. <https://doi.org/10.1177/1536867X1001000309>.
- Hampel, F. R. 1974. The influence curve and its role in robust estimation. *Journal of the American Statistical Association* 69: 383–393. <https://doi.org/10.2307/2285666>.
- Heckley, G., U.-G. Gerdtham, and G. Kjellsson. 2016. A general method for decomposing the causes of socioeconomic inequality in health. *Journal of Health Economics* 48: 89–106. <https://doi.org/10.1016/j.jhealeco.2016.03.006>.
- Jann, B. 2008. The Blinder–Oaxaca decomposition for linear regression models. *Stata Journal* 8: 453–479. <https://doi.org/10.1177/1536867X0800800401>.
- Kolenikov, S. 2010. Resampling variance estimation for complex survey data. *Stata Journal* 10: 165–199. <https://doi.org/10.1177/1536867X1001000201>.
- von Mises, R. 1947. On the asymptotic distribution of differentiable statistical functions. *Annals of Mathematical Statistics* 18: 309–348. <https://doi.org/10.1214/aoms/1177730385>.
- Oaxaca, R. 1973. Male–female wage differentials in urban labor markets. *International Economic Review* 14: 693–709. <https://doi.org/10.2307/2525981>.
- Rothe, C. 2010. Nonparametric estimation of distributional policy effects. *Journal of Econometrics* 155: 56–70. <https://doi.org/10.1016/j.jeconom.2009.09.001>.

About the author

Fernando Rios-Avila is a research scholar at Levy Economics Institute of Bard College under the Distribution of Income and Wealth program. His research interests include applied econometrics, labor economics, and poverty and inequality.

A Functional statistics and RIFs

This appendix provides the full set of distributional statistics, formulas, and sources.

Table A.1. Functional statistics and RIFs

Statistic	Definition	RIF	Source
Mean	$\mu_Y = \int y dF_Y(y)$	$\text{RIF}(y, \mu_Y) = y$	Firpo, Fortin, and Lemieux (2018)
Variance	$\sigma_Y^2 = \int (y - \mu_Y)^2 dF_Y(y)$	$\text{RIF}(y, \sigma_Y^2) = (y - \mu_Y)^2$	Firpo, Fortin, and Lemieux (2018)
p th quantile	$q_Y(p) = F_Y^{-1}(p)$	$\text{RIF}\{y, q_Y(p)\} = q_Y(p) + \frac{p - 1\{y \leq q_Y(p)\}}{f\{q_Y(p)\}}$	Firpo, Fortin, and Lemieux (2018)
Interquantile range	$\text{iqr}_Y(p_1, p_2)$ $= q_Y(p_2) - q_Y(p_1)$	$\text{RIF}\{y, \text{iqr}_Y(p_1, p_2)\} = \text{RIF}\{y, q_Y(p_1), F_Y\}$ $- \text{RIF}\{y, q_Y(p_2), F_Y\}$	Firpo, Fortin, and Lemieux (2018)
Gini	$\text{Gini}_Y = 1 - \frac{2}{\mu_Y} R_Y$ $R_Y = \int_0^1 \text{GL}_Y(p) dp$ $\text{GL}_Y(p) = \int_{-\infty}^{q_Y(p)} y dF_Y(y)$	$\text{RIF}(y, \text{Gini}_Y) = 1 + \frac{2}{\mu_Y^2} R_Y - \frac{2}{\mu_Y} [y\{1 - F_Y(y)\}]$	Firpo, Fortin, and Lemieux (2018)

Continued on next page

Statistic	Definition	RIF	Source
Coefficient of variation	$cv_Y = \frac{\sigma_Y}{\mu_Y}$	$RIF(y, cv_Y) = cv_Y + \frac{1}{2} \frac{(y - \mu_Y)^2 - \sigma_Y^2}{\mu_Y \times \sigma_Y} - \frac{\sigma_Y}{\mu_Y^2} (y - \mu_Y)$ $v = \int y^2 dF_Y(y)$	Firpo and Pinto (2016)
Standard deviation	$\sigma_Y = \sqrt{\int (y - \mu_Y)^2 dF_Y(y)}$	$RIF(y, \sigma_Y) = \sigma_Y + \frac{1}{2} \frac{(y - \mu_Y)^2 - \sigma_Y^2}{\sigma_Y}$	No source
Interquantile ratio	$iqratio_Y(p_1, p_2) = \frac{q_Y(p_2)}{q_Y(p_1)}$	$RIF\{y, iqratio_Y(p_1, p_2)\} = iqratio_Y(p_1, p_2)$ $+ \frac{1}{q_Y(p_1)} \left[\frac{p_2 - 1 \{y \leq q_Y(p_2)\}}{f\{q_Y(p_2)\}} \right]$ $- \frac{q_Y(p_2)}{q_Y(p_1)} \frac{p_1 - 1 \{y \leq q_Y(p_1)\}}{f\{q_Y(p_1)\}} \Big]$	Chung and Van Kerm (2018)
Generalized entropy index, $a \notin \{0, 1\}$	$I_{Y_E}^\alpha = \frac{1}{\alpha(\alpha - 1)} \left(\frac{v}{\mu_Y^\alpha} - 1 \right)$ $v = \int y^\alpha dF_Y(y)$	$RIF(y, I_{Y_E}^\alpha) = I_{Y_E}^\alpha + \frac{y^\alpha - v}{\alpha(\alpha - 1)\mu_Y^\alpha} - \frac{v}{(\alpha - 1)\mu_Y^{\alpha+1}} (y - \mu_Y)$	Cowell and Flachaire (2007)
Generalized entropy index, $a = 1$	$I_{Y_E}^1 = \frac{v}{\mu_Y} - \log \mu_Y$ $v = \int y \ln y dF_Y(y)$	$RIF(y, I_{Y_E}^1) = I_{Y_E}^1 + \frac{1}{\mu_Y} (y \ln y - v) - \frac{v + \mu_Y}{\mu_Y^2} (y - \mu_Y)$	Cowell and Flachaire (2007)
Generalized entropy index, $a = 0$	$I_{Y_E}^0 = \log \mu_Y - v$ $v = \int \ln y dF_Y(y)$	$RIF(y, I_{Y_E}^0) = I_{Y_E}^0 - (\ln y - v) + \frac{1}{\mu_Y} (y - \mu_Y)$	Cowell and Flachaire (2007)

Continued on next page

Statistic	Definition	RIF	Source
Atkinson index $\varepsilon > 0$ and $\varepsilon \neq 1$	$I_A^\varepsilon = 1 - \frac{v^{\frac{1}{1-\varepsilon}}}{\mu_Y}$ $v = \int y^{1-\varepsilon} dF_Y(y)$	$\text{RIF}(y, I_Y^\varepsilon) = I_A^\varepsilon + \frac{v^{\frac{\varepsilon}{1-\varepsilon}}}{(\varepsilon - 1)\mu_Y} (y^{1-\varepsilon} - v)$ $+ \frac{v^{\frac{1}{1-\varepsilon}}}{\mu_Y^2} (y - \mu_Y)$	Cowell and Flachaire (2007)
Atkinson index $\varepsilon = 1$	$I_A^1 = 1 - \frac{e^v}{\mu_Y}$ $v = \int \ln y dF_Y(y)$	$\text{RIF}(y, I_Y^1) = I_A^1 - \frac{e^v}{\mu_Y} (\ln y - v) + \frac{e^v}{\mu_Y^2} (y - \mu_Y)$	Cowell and Flachaire (2007)
Logarithmic variance	$\text{LV}_Y = \int \left(\ln \frac{y}{\mu_Y} \right)^2 dF_Y(y)$	$\text{RIF}(y, \text{LV}_Y) = \text{LV}_Y + \{(\log y)^2 - v_1\}$ $- 2\log \mu_Y (\ln y - v_2)$ $- \frac{2}{\mu_Y} (v_2 - \ln \mu_Y)(y - \mu_Y)$ $v_1 = \int (\ln y)^2 dF_Y(y); \quad v_2 = \int y dF_Y(y)$	Cowell and Flachaire (2007)
Generalized Lorenz ordinate	$\text{GL}_y(p) = \int_{-\infty}^{q_Y(p)} y dF_Y(y)$	$\text{RIF}\{y, \text{GL}_Y(p)\} = pq_Y(p) + \{y - q_Y(p)\}\{y < q_Y(p)\}$	Essama- Nssah and Lambert (2012)

Continued on next page

Statistic	Definition	RIF	Source
Lorenz ordinate	$L_Y(p) = \frac{GL_y(p)}{\mu_Y}$	$\text{IF}\{y, L_Y(p)\} = -\frac{y}{\mu_Y} L_Y(p) + \frac{pq_Y(p)}{\mu_Y}$ $+ \left\{ \frac{y - q_Y(p)}{\mu_Y} \right\} \{y < q_Y(p)\}$ $\text{RIF}\{y, L_Y(p)\} = L_Y(p) + \text{IF}\{y, L_Y(p)\}$	Essama-Nssah and Lambert (2012)
Upper class share	$\text{ucs}_Y(p) = 1 - L_Y(p)$	$\text{RIF}\{y, \text{ucs}_Y(p)\} = \text{ucs}_Y(p) - \text{IF}\{y, L_Y(p)\}$	No source
Interquantile share ratio	$\text{Iqrs}_Y(p_1, p_2) = \frac{1 - L_Y(p_2)}{L_Y(p_1)}$	$\text{RIF}\{y, \text{Iqrs}_Y(p_1, p_2)\}$ $= \text{Iqrs}_Y + \frac{1}{L_Y(p_1)}$ $\left[-\text{IF}\{y, L_Y(p_2)\} - \frac{1 - L_Y(p_2)}{L_Y(p_1)} \text{IF}\{y, L_Y(p_1)\} \right]$	No source
Middle class share	$\text{mcs}_Y(p_1, p_2) = L_Y(p_2) - L_Y(p_1)$	$\text{RIF}\{y, \text{mcs}_Y(p_1, p_2)\} = \text{RIF}\{y, L_Y(p_2)\} - \text{RIF}\{y, L_Y(p_1)\}$	Davies, Fortin, and Lemieux (2017)
FGT poverty indices	$\text{FGT}_Y(\alpha, Z)$ $= \int_{-\infty}^Z \left(\frac{Z - y}{Z} \right)^\alpha dF_Y(y)$ $\alpha \geq 0, Z = \text{poverty line}$	$\text{RIF}\{y, \text{FGT}_Y(\alpha, Z)\} = \left(\frac{Z - y}{Z} \right)^\alpha (y \leq Z)$	Essama-Nssah and Lambert (2012)

Continued on next page

Statistic	Definition	RIF	Source
Watts index	$W_Y(Z) = \int_0^Z \ln \frac{Z}{Y} dF_Y(y)$	$\text{RIF}\{y, W_Y(Z)\} = \ln \frac{Z}{y} (y < Z)$	Essama-Nssah and Lambert (2012)
Sen index	$S_Y(Z) = \frac{2}{ZF_Y(Z)} \int_0^Z (z - y) \{F_Y(Z) - F_Y(y)\} dF_Y(y)$	$\text{RIF}\{y, S_Y(Z)\} = -\frac{1}{F_Y(Z)} S_Y(Z) - \frac{2}{ZF_Y(Z)} \int_0^y \{F_Y(Z) - F_Y(x)\} dx + 2$	Essama-Nssah and Lambert (2012)
TIP curve ordinate	$\text{TIP}_Y(Z, p) = \int_0^x (z - y) dF_Y(y)$ $x = \min\{Z, q_Y(p)\}$	$\text{RIF}\{y, \text{TIP}_Y(Z, p)\} = \begin{cases} (Z - y) \times (Z > y) & \text{if } Z < q_Y(p) \\ p\{Z - q_Y(p)\} + \{q_Y(p) - y\}\{q_Y(p) < y\} & \text{if } Z \geq q_Y(p) \end{cases}$	Essama-Nssah and Lambert (2012)
Absolute Gini	$\text{Agin}_Y = 2 \int_{-\infty}^{\infty} (y - \mu_Y) \{F_Y(y) - 0.5\} dF_Y(y)$ $= 2\text{Cov}\{y, F_Y(y)\}$	$\text{RIF}(y, \text{Agin}_Y) = -\text{Agin}_Y + (\mu_Y - y) 2[yF_Y(y) - \text{GL}_Y\{F_Y(y)\}]$	Essama-Nssah and Lambert (2012)

For the following indices, one assumes that the data used can be written as $(H, Y) = [(h_1, y_1), (h_2, y_2), \dots, (h_n, y_n)]$. The joint probability functions for H and F_Y are f_{H, F_Y} and F_{H, F_Y} , and the data contamination is $\delta_{h,y}(h_c, y_c) = 1$ if $h_i \geq h_c$ and $F(y_1) \geq F(y_c)$.

Absolute concentration index	$\text{ACI}(h, F_{H, F_Y}) = 2\text{Cov}\{h, F_Y(y)\}$ where h is variable of interest and y is ranking variable	$\text{RIF}\{h, \text{ACI}(h, F_{H, F_Y})\} = \text{ACI}(h, F_{H, F_Y}) + \text{IF}\{h, \text{ACI}(h, F_{H, F_Y})\}$ $\text{IF}\{h, \text{ACI}(h, F_{H, F_Y})\} = -2\text{ACI}(h, F_{H, F_Y}) + (\mu_H - h) + 2\left\{hF_{H, F_Y} - \int^y \int^\infty h f_{H, F_Y} dh dF_Y(x)\right\}$	Heckley, Gerdtham, and Kjellsson (2016)
------------------------------	---	--	---

Continued on next page

Statistic	Definition	RIF	Source
Concentration index	$CI(h, F_Y) = \frac{ACI(h, F_{H,F_Y})}{\mu_H}$	$RIF\{h, CI(h, F_{H,F_Y})\} = CI(h, F_{H,F_Y}) + \frac{\mu_H - h}{\mu_H^2} ACI(h, F_{H,F_Y}) + \frac{1}{\mu_H} IF\{h, ACI(h, F_{H,F_Y})\}$	Heckley, Gerdtham, and Kjellsson (2016)
Erreygers index	$EI(h, F_{H,F_Y}, ub, lb) = \frac{4ACI(h, F_{H,F_Y})}{ub - lb}$	$RIF\{h, EI(h, F_{H,F_Y}, ub, lb)\} = EI(h, F_{H,F_Y}, ub, lb) + \frac{4}{ub - lb} IF\{h, ACI(h, F_{H,F_Y})\}$	Heckley, Gerdtham, and Kjellsson (2016)
Attainment relative concentration index	$ARI(h, F_{H,F_Y}, lb) = \frac{ACI(h, F_{H,F_Y})}{\mu_H - lb}$	$RIF\{h, ARI(h, F_{H,F_Y}, lb)\} = ARI(h, F_{H,F_Y}, lb) + \frac{\mu_H - h}{(\mu_H - lb)^2} ACI(h, F_{H,F_Y}) + \frac{1}{\mu_H - lb} IF\{h, ACI(h, F_{H,F_Y})\}$	Heckley, Gerdtham, and Kjellsson (2016)
Shortfall relative concentration index	$SRI(h, F_{H,F_Y}, ub) = \frac{ACI(h, F_{H,F_Y})}{ub - \mu_H}$	$RIF(h, SRI\{h, F_{H,F_Y}, ub\}) = SRI(h, F_{H,F_Y}, ub) + \frac{h - \mu_H}{(ub - \mu_H)^2} ACI(h, F_{H,F_Y}) + \frac{1}{ub - \mu_H} IF\{h, ACI(h, F_{H,F_Y})\}$	Heckley, Gerdtham, and Kjellsson (2016)

Continued on next page

Statistic	Definition	RIF	Source
Wagstaff index	$\text{WI}(h, F_{H, F_Y}, \text{ub}, \text{lb})$ $= \frac{(\text{ub} - \text{lb})\text{ACI}(h, F_{H, F_Y})}{(\text{ub} - \mu_H)(\mu_H - \text{lb})}$	$\text{RIF}\{h, \text{WI}(h, F_{H, F_Y}, \text{ub}, \text{lb})\}$ $= \text{WI}(h, F_{H, F_Y}, \text{ub}, \text{lb})$ $+ \frac{(h - \mu_H)(\text{ub} - \text{lb})(\text{ub} + \text{lb} - 2\mu_H)}{\{(\text{ub} - \mu_H)(\mu_H - \text{lb})\}^2} \text{ACI}(h, F_{H, F_Y})$ $+ \frac{\text{ub} - \text{lb}}{(\text{ub} - \mu_H)(\mu_H - \text{lb})} \text{IF}\{h, \text{ACI}(h, F_{H, F_Y})\}$	Heckley, Gerdtham, and Kjellsson (2016)

B RIF statistics and statistical inference

This appendix provides the Monte Carlo simulation results to evaluate the use of RIFs for statistical inference regarding distributional statistics. This particular use of RIFs has been discussed in Cowell and Flachaire (2015), Deville (1999), and Efron (1982).

For the statistics provided below, I use a sample of 2,500 observations of two variables, \mathbf{x}_1 and \mathbf{x}_2 , drawn from a jointly standard normal distribution, with variance 1 and correlation = 0.5. To simulate data that resemble income distributions, the random draws for \mathbf{x}_1 and \mathbf{x}_2 are used to create draws from a chi-squared distribution with 5 degrees of freedom, using an inverse-transformation approach:

$$z_{i,k} = F_{\chi^2(5)}^{-1}\{\Phi(x_{i,k})\} \quad \text{for } k = 1 \text{ and } 2$$

$F_{\chi^2(5)}^{-1}(\cdot)$ is the inverse cumulative function corresponding to a chi-squared distribution with 5 degrees of freedom, and $\Phi(\cdot)$ is the normal c.d.f. Based on this data structure, 10,000 repetitions are drawn and the RIFs detailed in table 7 are obtained to estimate the RIF's standard errors. For the bivariate distributional concentration, indices are estimated for variable $z_{i,1}$ based on the ranking from $z_{i,2}$.

A simple look at the results, particularly the ratio between the simulated standard error and the average standard error obtained from using the RIFs, shows that for most statistics the results are robust, with two exceptions.

First, the largest biases seem to be associated with the estimation of standard errors for sample quantile, interquantile range, and interquantile ratio statistics, in particular when using quantiles at the lower end of the distribution (10th). On average, the direct use of RIFs for estimating standard errors overstates by almost 10%, compared with the simulated standard errors. This has also been described in Firpo, Fortin, and Lemieux (2009), who indicate that the estimation of the sample density increases the complexity for the estimation of the asymptotic standard errors for UQRs, suggesting instead the use of bootstrap standard errors.

Second, one also observes that the variance associated with the Atkinson statistics understates the statistic standard errors by almost 9% for an inequality aversion = 2. Additional simulations not shown here suggest that the size of the bias increases with the degree of inequality aversion but is negligible for low levels of inequality aversion. This also suggests the use of bootstrap standard errors when one is interested in drawing conclusions regarding this inequality index.

The use of RIFs for the estimation of standard errors seems to be robust for all other statistics, with an average bias of less than 1% based on the current simulation.

Table B.2. Simulation results: Evaluating asymptotic performance of RIFs for the estimation of statistics sample errors

Statistic	Distributional statistic	Simulation- based standard error	Avg. RIF standard error	Ratio ¹
Mean	5.0003	0.0630	0.0632	1.0031
Variance	9.9946	0.4170	0.4166	0.9991
10th quantile	1.6111	0.0490	0.0533	1.0865
50th quantile	4.3525	0.0727	0.0739	1.0160
90th quantile	9.2384	0.1623	0.1596	0.9831
50–10 interquantile range	2.7414	0.0733	0.0753	1.0276
90–50 interquantile range	4.8859	0.1538	0.1520	0.9877
Gini index	0.3394	0.0044	0.0044	1.0030
Coefficient of variation	0.6321	0.0106	0.0105	0.9945
Standard deviation	3.1607	0.0659	0.0658	0.9989
50/10 interquantile ratio	2.7037	0.0801	0.0860	1.0740
90/50 interquantile ratio	2.1229	0.0420	0.0420	1.0019
Entropy index $e = 0$	0.2130	0.0060	0.0060	1.0006
Entropy index $e = 1$	0.1868	0.0050	0.0050	1.0006
Entropy index $e = 2$	0.1998	0.0067	0.0066	0.9944
Atkinson index $a = 1$	0.1919	0.0048	0.0049	1.0007
Atkinson index $a = 1.5$	0.2930	0.0079	0.0079	0.9933
Atkinson index $a = 2$	0.3995	0.0144	0.0132	0.9140
Logarithmic variance	0.5355	0.0192	0.0192	0.9972
Generalized Lorenz ordinate at $p = 20$	0.3079	0.0080	0.0080	1.0065
Generalized Lorenz ordinate at $p = 40$	0.9080	0.0174	0.0175	1.0060
Generalized Lorenz ordinate at $p = 60$	1.7812	0.0285	0.0286	1.0052
Generalized Lorenz ordinate at $p = 80$	3.0037	0.0423	0.0423	1.0011
Lorenz ordinate at $p = 20$	0.0616	0.0014	0.0014	1.0042
Lorenz ordinate at $p = 50$	0.2616	0.0031	0.0031	1.0051
Lorenz ordinate at $p = 80$	0.6007	0.0037	0.0037	1.0035
Upper class share at $p = 20$	0.9384	0.0014	0.0014	1.0042
Upper class share at $p = 50$	0.7384	0.0031	0.0031	1.0051
Upper class share at $p = 80$	0.3993	0.0037	0.0037	1.0035
Interquantile share ratio 90–10	10.8464	0.4258	0.4252	0.9987
Interquantile share ratio 80–20	6.4894	0.1851	0.1856	1.0028
Interquantile share ratio 60–40	3.5463	0.0692	0.0694	1.0033
Middle class share 10–90	0.7422	0.0031	0.0031	1.0049
Middle class share 20–80	0.5391	0.0033	0.0033	1.0053
Middle class share 40–60	0.1746	0.0016	0.0016	1.0041
FGT($a = 0, Z = 2.5$) headcount	0.2235	0.0084	0.0083	0.9949
FGT($a = 1, Z = 2.5$) poverty gap	0.0777	0.0036	0.0036	1.0058
FGT($a = 2, Z = 2.5$) poverty severity	0.0389	0.0023	0.0023	1.0082
Watts poverty index $Z = 2.5$	0.1154	0.0062	0.0062	1.0061
Sen poverty index $Z = 2.5$	0.1072	0.0047	0.0048	1.0046
TIP ordinate at $p = 10z = 2.5$	0.1411	0.0039	0.0039	1.0034
TIP ordinate at $p = 25z = 2.5$	0.1942	0.0090	0.0091	1.0058
TIP ordinate at $p = 50z = 2.5$	0.1942	0.0090	0.0091	1.0058
Absolute Gini	1.6963	0.0307	0.0308	1.0028
Absolute concentration index	0.8521	0.0356	0.0354	0.9948
Concentration index	0.1705	0.0066	0.0066	0.9941
Erreygers's index $lb(1) \ ub(9)$	0.4261	0.0178	0.0177	0.9948
Attainment relative concentration				

Continued on next page

Statistic	Distributional statistic	Simulation-based standard error	Avg. RIF standard error	Ratio ¹
index lb(1)	0.2130	0.0082	0.0082	0.9947
Shortfall relative concentration				
index ub(9)	0.2132	0.0106	0.0106	0.9965
Wagstaff index lb(1) ub(9)	0.4262	0.0178	0.0177	0.9948

NOTE: Monte Carlo simulation using 10,000 repetitions.

¹ Ratio is defined as the ratio between the average RIF standard errors and the simulation-based standard errors.

C Additional commands for the estimation of UQRs and simultaneous RIF regressions

C.1 uqreg

The most popular application of RIF regressions in the literature is the estimation of UQRs. As described in the main text, the command `rifhdreg` can be used to estimate RIF regressions for a large set of distributional statistics, including quantile effects. A disadvantage is that the RIF regressions are estimated using OLS. Firpo, Fortin, and Lemieux (2009) explain that for the particular case of unconditional quantiles, one can also use binomial models, such as logit or probit, for the estimation of the UQR.

Recall that the RIF corresponding to quantiles is defined as

$$\text{RIF}\{y, q_Y(p)\} = q_Y(p) + \frac{p - 1\{y \leq q_Y(p)\}}{f\{q_Y(p)\}}$$

After some rearrangements, this can be written as

$$\text{RIF}\{y, q_Y(p)\} = q_Y(p) + \frac{p - 1}{f\{q_Y(p)\}} + \frac{1\{y > q_Y(p)\}}{f\{q_Y(p)\}}$$

This equation implies that once p is determined and the quantile $q_Y(p)$ and density of the distribution at that quantile, $f\{q_Y(p)\}$, have been estimated, the only parameter that needs modeling is $1\{y > q_Y(p)\}$, which is a binary variable. In other words, when we fit the RIF regression model using OLS, one is essentially using a linear probability model to analyze this variable. Because of this, Firpo, Fortin, and Lemieux (2009) suggest that one can also use binomial models, such as logit or probit models, for better model specification. In this case, the UPE could be estimated using the rescaled average marginal effects.

Consider a more general specification of the estimation of UQR:

$$\text{RIF}\{y, q_Y(p)\} = q_Y(p) + \frac{p - 1}{f\{q_Y(p)\}} + \frac{H(\mathbf{X}\boldsymbol{\gamma})}{f\{q_Y(p)\}}$$

where $H(\mathbf{X}\boldsymbol{\gamma})$ is some linear or nonlinear function used to model $1\{y > q_Y(p)\}$. In this case, the UPE can be estimated as follows:

$$E \left[\frac{\partial \text{RIF}\{y, q_Y(p)\}}{\partial X} \right] = \frac{1}{f\{q_Y(p)\}} E \left\{ \frac{\partial H(\mathbf{X}\boldsymbol{\beta})}{\partial X} \right\}$$

For the estimation of UQR using other model specifications, I suggest using the command `uqreg`. This standalone command models a binary variable D defined as $1\{y > q_Y(p)\}$ using other alternatives chosen by the user. The syntax is as follows:

```
uqreg depvar [indepvars] [if] [in] [weight], q(#p) method(str) [bw(#)  
      kernel(kernel) method_options noisily]
```

aweights, **fweight**s, **ilweight**s, and **pweight**s can be used depending on the methods chosen; see [U] **11.1.6 weight**.

q(*#p*) indicates the quantile of interest, with $0 < \#p < 100$. **q**() is required.

method(*str*) indicates the method that will be used to model $1\{y > q_Y(p)\}$. Any method can potentially be specified here, as long as average marginal effects can be estimated using **margins**. Examples using **regress**, **logit**, **probit**, and **cloglog** are provided in the help file. **method**() is required.

bw(*#*) and **kernel**(*kernel*) specify the bandwidth, kernel function, or both that will be used to estimate the density function. The defaults are **kernel**(**gaussian**) and the Silverman plugin bandwidth.

method_options are any options appropriate for the method selected.

noisily displays preliminary steps.

The output of this command reports UPE for all variables used in the model, but it will not report standard errors. Standard errors can be obtained using bootstrap methods.

C.2 rifsureg and rifsureg2

Another disadvantage of **rifhdreg** is that one can estimate the impact of independent variables for only one distributional statistic at a time. However, researchers are often interested in testing if the impact of independent variables varies across quantiles or may be interested in analyzing multiple statistics simultaneously. One solution to address this type of problem, in a more general framework, is to use the Stata command **suest**. Unfortunately, the estimation output from **rifhdreg** is not compatible with **suest**.

An alternative method with a narrower scope is the estimation of simultaneous linear regression models using the command `sureg`. This command estimates seemingly unrelated regressions when one believes the errors across models are correlated. Because RIF regressions are estimated via OLS, `sureg` can be adapted to fit simultaneous models.³³

In this context, `rifsureg` and `rifsureg2` are wrappers for `sureg` that can be used to fit simultaneous RIF regression models. `rifsureg` concentrates on the estimation of simultaneous UQR and can be considered equivalent to `sqreg`, which fits simultaneous conditional quantile regression models. `rifsureg2` can be used more broadly for all other statistics.

The syntaxes of `rifsureg` and `rifsureg2` are similar to that for `rifhdreg`:

```
rifsureg depvar [indepvars] [if] [in] [weight], qs(numlist) [bw(#)]
        kernel(kernel) retain(str) replace over(varname) rwlogit(varlist)
        rwprobit(varlist) rwmlogit(varlist) rwmprobit(varlist) [ate|att|atu]
        sureg_options]

rifsureg2 depvar [indepvars] [if] [in] [weight], rif(RIF_options)
        [retain(str) replace over(varname) rwlogit(varlist) rwprobit(varlist)
        rwmlogit(varlist) rwmprobit(varlist) [ate|att|atu] sureg_options]
```

`fweights` and `awweights` are allowed. When using `rwlogit()`, `rwprobit()`, `rwmlogit()`, or `rwmprobit()`, weights are used as `awweights`; see [U] 11.1.6 **weight**.

The main difference from `rifhdreg` is that neither of these commands can be used for the estimation of RIF regressions with fixed effects. `rifsureg` requests the estimation of UQR, fitting all the models declared in `qs(numlist)`. The list of quantiles can be any whole number from 1 to 99. The selection of the kernel function and bandwidth can be done directly as options of the command. `rifsureg2` requests the estimation of all the statistics that are indicated with the option `rif()`, separating each statistic with a comma. Each statistic follows the same format as in `rifhdreg`. For example, if one would like to estimate RIF regressions for the 10th quantile, the variance, and the Gini, one could type `rifsureg2 y x1 x2, rif(q(10), var, gini)`.

After `rifsureg` is used, new variables will be automatically created containing the quantile RIFs for all quantiles declared in `qs(numlist)` under the name `__varname_q#`. After `rifsureg2` is used, new variables will be created under the name `__varname_m#`, where the number indicates the order of the statistic in the list.

While the command reports standard errors estimated by `sureg`, one should instead estimate bootstrap standard errors with the prefix `bootstrap`.

33. In the last example in the help file for `rifhdreg`, I provide an alternative example of how one can fit simultaneous RIF regression models using the command `suest`.